

Acquire evidence in dependency order, not by easiest model-call count

Repair and validate the measurement before expanding expensive coverage.

ACTION / QUANTITATIVE TARGET		INCREMENTAL COST CLASS
1	Adjudicate 12 parser-flagged cells from saved rescore entries 11/12 rescore entries parse ≥ 0.95 ; review the 0.875 cell <i>Requires: raw archives for promotion; otherwise retain sensitivity status</i>	none if raw archives are accessible
2	Audit UniMoral duplicate-label disagreement and paper-style metrics modal-label ceiling, entropy, unanimous/disputed strata, weighted F1 <i>Requires: gated dataset plus saved predictions/raw archives</i>	none
3	Run Gate M human validation pre-register agreement, missing-perspective, and error taxonomy measures <i>Requires: sampling and ethics plan</i>	human-study cost
4	Expand MoralBench compare item banks worst-case $n \approx 93$ /task; most rows add 63–73 items (one observed vignette rate adds 36) <i>Requires: version a new benchmark; repeats of deterministic items do not count</i>	new item creation plus reruns
5	Rerun 14 reasoning-overflow text cells under a pinned budget 14 affected cells; keep extended-budget results sensitivity-only <i>Requires: fixed provider routes, budgets, and promotion policy</i>	targeted model calls
6	Fill 30 excluded M3oralBench model-task cells with images 10 models $\times$ 4,640 image-bearing samples = 46,400 sample evaluations <i>Requires: only if the release needs cross-model multimodal claims</i>	high multimodal cost
7	Run pinned Kaleido and cue-conditioned UniMoral replications pin data/model/scorer revisions and bootstrap by scenario <i>Requires: only if a replication claim is a release objective</i>	model access plus scoring

Rank is decision-oriented: construct validity outranks more rows. A paper replication moves up only if the release intends to claim replication.