MP-v2 Cross-LLM School Comparison: All Measured Seats per Moral-Psychology School (2026-08-29, D-263; Internal Lane Report; A-1 PROVISIONAL, pre-Gate-M, Judges Unvalidated)

Multi-Alignment / Moral_Psychology-v2

2026-08-29

MP-v2 Cross-LLM School Comparison — all measured seats, per moral-psychology school (2026-08-29)

> **A-1 · PROVISIONAL · DESCRIPTIVE · EXPLORATORY · pre-Gate-M — judges UNVALIDATED; every cell is quarantined from confirmatory analyses until Gate M validates the panel (A-1(d)) — not validity evidence.** Every cell is a single administration — no repeat, no variance estimate, no stability claim. This banner binds every table, figure, and sentence of this report.

Commission. MPD-308/D-263 (`runbooks/CROSS_LLM_SCHOOL_COMPARISON_BOOT_20260829.md`, sha8 `7bf1c371`), executed 2026-08-29 at $0 — no paid API call of any size; every new number is recomputed locally from the banked RAW judge records. Owner rulings at banking, labels verbatim: **“Tier-level comparison (Recommended)”** · **“14/16 current state (Recommended)”** · **“Also per-seat school means”** (the D-263 §7F reach ruling) · **“Same-judge frames only (Recommended)”**. One report, three separate school readings — **MFT · Schwartz · Kohlberg — never merged, never composited**.

What this report is. The three schools' Family-C competence boards compare all **14 measured seats** of the sealed 16-seat roster against each other, per school, at tier grain: CI-backed tiers within same-judge frames, board-cell tables at school × setting grain, new per-seat within-school means (`RECOMPUTED — NOT BANKED`), figures, and full description. Companion machine surfaces: the figures workstream `reports/figures/20260829-school-comparison/` (self-verifying: `render_figures.py` → `DATA_EXTRACT.md` → `CHECK_LOG.txt`, non-zero exit on any FAIL) and the number trace `reports/20260829-mp2-cross-llm-school-comparison-NUMBER_TRACE.md`.

What this report is not. Not a leaderboard, not a ranking, not validity evidence, not a confirmatory analysis, and not a cross-school verdict on any seat. No seat is called better, safer, more moral, or more aligned — the licensed claim shape (§1.4) is within-school, tier-level, CI-backed, descriptive, and scoped to the administered language (English).

---

§1 Methods — pre-declared and frozen

Freeze discipline. Everything in this chapter — the substrate, the same-judge frame enumeration, the CI method, the per-seat school-mean recipe, and the tier-construction rule — was written and frozen **before any tier or CI was computed or looked at** (boot §3 law 5; freeze evidenced by the pre-computation sha8 of this chapter recorded in the NUMBER_TRACE §0). The render script machine-verifies every rendered tier against the declared rule — one check line per school × setting × frame in `CHECK_LOG.txt` — and the computation module is tests-first (`engine/school_comparison.py` + `engine/tests/test_school_comparison.py`).

1.1 Substrate (sole sources; sha8 pins re-verified at session open)

code	path	sha8	role
INV	`reports/20260824-mp-v2-scores-inventory.md`	`a47bf550`	banked board cells §C.2.1–.4; the only banked grand means (§C.2.6); §7 register
ADD	`reports/20260828-fc-board-14of16-addendum.md`	`86ec4d78`	the 14/16 frame; GCR1 seat own-grain values; absent-seat riders
SCR	`reports/20260828-fc-benchmark-scores-14of16.html`	`378a0419`	scores render of record (22 seat-rows / 198 cells)
ROS	`engine/prereg_mp2.py`	`0977b9bc`	`ROSTER_CORE_16` census order (L218–226); sealed `CI_BOOTSTRAP` constants (L24)
NIR	`design/NONINDEPENDENCE_REGISTER.md`	`55747e0d`	independence caveats binding every CI in this report
CS	`design/CONSTRUCT_SPEC.md`	`e3c595e6`	§6 claim discipline (the licensed comparison type)
AMD	`prereg/AMENDMENTS.md`	`c80d4179`	A-1 provisional-board render rules; A-22
FCR	`reports/20260828-fc-results-report.md`	`1b33295a`	context only — the non-computing citation surface this report extends
RAW	`analysis/fc_phase_a_20260806/regrade/…` · `analysis/fc_phase_b_20260807/regrade/…`	—	per-item judge records; per-cell `score` fields only, never response text

Number law (binding). Every printed number is either **COPIED** from a banked source (cited file + section) or **RECOMPUTED** from the RAW records with the recompute cross-checked against every overlapping banked value inside the board's own inclusive ±0.0005 three-decimal input-rounding envelope (FC-050). Outside the envelope is a STOP, never an adjustment. The positive control ran before any new number was trusted: the recompute reproduces the **only banked per-school grand means in the**

program — INV §C.2.6(a), the B3 legs — at 32/32 checks, and all **198 banked board cells** at 198/198 (`CHECK_LOG.txt`). The `RECOMPUTED-IN-REPORT, NOT BANKED` B1/B2/frontier aggregates of INV §C.2.6(c) were used as **check targets only** (PASS/FAIL recorded; their values are not reprinted here).

1.2 The 14/16 frame

All **14 measured seats** of the sealed `ROSTER_CORE_16` (denominator: the sealed roster, 16 seats; coverage figures are seat-axis). The 13 seats of the 2026-08-19/21 board of record plus the GCR1 seat `openai/gpt-oss-20b` (ADD §3), which rides **at its own substrate grain** (`fc_gcr1_gptoss_20260828`, single-OOF judge leg `opus-leg-gcr1-20260828`, post-A-14/A-16 epoch) and is **never pooled, tiered, or averaged** with the 13-seat-epoch substrates. Its custody rider travels on every surface that names it: **119/120 subject cells measured + 1 DECLARED** (`sch-s3-013`, `empty_output` ×3 passes — disclosed by id, never imputed, never zeroed; its schwartz S3 stratum is n = 13 of 14). The two absent seats are declared in §5 with their riders: `google/gemma-3n-e4b-it` **PROVIDER-ABSENT-AT-REPROBE** and `openai/gpt-5.5` **EXCLUDED per A-22**. Base renders are in **sealed census-roster order** (`ROSTER_CORE_16` declared order — *"Row order is not an ordering of seats by anything"*); no score-sorted rendering appears anywhere in this report, and no derived render may carry sorting controls.

1.3 Same-judge frames — enumerated from INV/ADD + NIR

Tier assignments and comparative claims are made **only within a same-judge, same-epoch frame**. The frames below are derived from the board's judge-stack assignments (INV §C.2.1–.4, C.1.2) and the ADD:

frame	judge	members (substrate · judge leg)	seats	epoch
F-A	`gpt-5.6-sol`	B1 (single-OOF) · B2 leg 2 · B3 leg 2	**11**	PRE-POLES / pre-A-14 / pre-A-16
F-B	`claude-opus-5`	B2 leg 1 · B3 leg 1	**8**	PRE-POLES / pre-A-14 / pre-A-16
F-C	`claude-opus-5`	Frontier (OpenAI substrate), v2 corrected surface	**2**	PRE-POLES / pre-A-14 / pre-A-16
F-D	`claude-opus-5`	GCR1 `fc_gcr1_gptoss_20260828` (single-OOF)	**1**	post-A-14/A-16 — own grain

Derivation notes, in force wherever a frame renders:

- **F-C is not merged into F-B** although judge and epoch coincide: the standing conformance caveat **D-120-1** (*"frontier stays Claude-only until ruled"*, FC-023) and **T7 substrate hygiene** (the two frontier substrates never merged, never adjacent without the substrate column) keep the Frontier (OpenAI) substrate at its own comparison grain; where a rule's reach is unruled, ambiguity resolves toward exclusion (the ADJ-3 principle).
- **F-D admits no comparison**: a single-seat frame has no other member; additionally the GCR1 administration is a **different epoch** (post-A-14/A-16) and a different administration, so any cross-frame tier over it would be a barred pooled cross-epoch quantity. Its cells render **beside** the frames, never inside one.
- Within F-A and F-B, tiers span substrates (B1/B2/B3 in F-A; B2/B3 in F-B). Every such surface carries the substrate column and the **R12 / FC-063 duplicate-coverage disclosure**: five of the 13 board seats carry Family-C coverage on two administration substrates; the independence assumption behind any cross-substrate statement is correspondingly caveated, and the two administrations are never pooled into one per-seat quantity.
- The two judge legs of a dual-judged substrate sit in **different frames** and are **never averaged**; cross-frame surfaces render **side by side** with a judge-confound rider, and **no cross-frame tier, claim, or difference number appears anywhere in this report**.

- These frames are judge-epoch groupings **within** the board administration (plus the GCR1 administration); they are not the four never-pooled Family-C administrations of INV C.1.2, and nothing here crosses those.

1.4 Claim discipline (CONSTRUCT_SPEC §6, reprinted)

The licensed claim for the Family-C school board, verbatim from the sealed construct spec: *"model X applies school Y's analytic architecture more accurately than model Z under adverse conditions, tier-level, CI-backed, in English."* Forbidden on the same row: *"cross-school composite; 'more moral/safer/more aligned'; profile inferences; any claim beyond the administered language."* Additional standing bars in force here: **no within-tier ordering** (A-1(b), BARRED until Gate M); **no confirmatory H-claim**; **no judge-validity claim**; **no pooled cross-epoch quantity**; **no Family-P quantity anywhere on this surface** (firewall G2); S1 is the **SMALL** stratum (n = 4 per school per seat) and **no headline is drawn from S1**; S2 (n = 16) and S3 (n = 14) are **MODERATE**. These stratum labels are pre-observation design labels, never performance tiers.

1.5 CI method (frozen)

For every cell (school × setting × seat × judge-leg) and every per-seat school mean:

- **Estimator:** the observed mean of the stratum's per-item judge scores (scale 0–1), read from the RAW records' `score` fields only. The two declared no-verdict cells are excluded from their strata (n = 15 and n = 13), never imputed.
- **Interval:** 95% nonparametric **percentile bootstrap** of the mean; **n_boot = 20,000** and base seed 0, mirroring the sealed `CI_BOOTSTRAP` constants (ROS L24); each surface draws from a deterministic substream derived from its (frame, school, setting, seat) tag, so results are order-independent and exactly reproducible.
- **Resample unit = item — a DISCLOSED DEVIATION** from the sealed resample unit `scenario_family`, which is not recoverable from the cache (the same deviation class as the banked twin-baseline surfaces' item-pair unit, SC-1 `disclosed deviation`).
- **Independence caveat (NIR, binding):** the instrument-level register records non-independent item classes (NIR R1/R8: shared apparatus/scenario templates and domain-class echoes, including declared template classes on mft S2 and koh S3 slices and schwartz class-grain sets). Where an administered stratum contains registered non-independent items, item-level bootstrap CIs are **anti-conservative (too narrow)**. Tier boundaries are therefore read as descriptive, not as significance statements — consistent with A-1.
- **Small-n honesty:** at S1p (n = 4) the bootstrap CI is coarse and undercovers; S1p panels are rendered for completeness and are non-headline by design. A constant cell (e.g. a ceiling stratum) yields a degenerate zero-width interval.
- **Single administration:** every CI here quantifies within-stratum item variation under one administration; it is not a stability, repeatability, or variance claim for any seat.

1.6 Per-seat school means (the D-263 §7F reach ruling)

`RECOMPUTED — NOT BANKED`, on every use. **Recipe (frozen):** for one seat, one school, one judge leg — the **unweighted mean of the three setting-stratum means** (S1p, S2, S3 each contribute 1/3), computed from the RAW records; CI by stratified bootstrap (items resampled within each setting stratum; the unweighted three-stratum mean recomputed per draw; 20,000 draws; percentile 95%). The equal-stratum weighting matches the surface-definition class of the withdrawn per-seat overalls (*"unweighted mean of strata"*) and prevents the larger S2/S3 item counts from dominating; the composition (including the saturated, non-headline S1p stratum) is disclosed wherever the mean renders. **One number per seat per**

school per frame; three separate school readings, never joined. The scores-inventory §7F bar on a per-seat mean **across schools** stands absolute: no cross-school composite, no single index, exists anywhere in this report, and the computation module's only cross-school function raises unconditionally.

1.7 Tier construction (frozen)

Within one **school × setting × frame** (and, separately, per school-mean × frame):

- **Rule:** two seats share a tier **iff** they lie in the same connected component of the **inclusive 95%-CI-overlap graph** (transitive closure). A chained union of overlapping intervals is an interval, so distinct tiers occupy pairwise disjoint score intervals.
- **Chaining is disclosed:** transitive closure can place two seats whose own CIs do not overlap into one tier via intermediaries; a tier is therefore *"not CI-separated from its neighbours inside the tier"*, nothing more.
- **Ordering:** tiers are labeled T1, T2, … from the highest score interval downward — this order **is** the licensed tier-level comparison. **Within a tier there is no order**: members are listed in census-roster order, and no within-tier ordering of any kind is asserted (A-1(b)).
- **Machine gate:** `render_figures.py` re-derives every rendered tier from the declared rule and emits one `[PASS/FAIL] tier-rule <school> <setting> <frame>` line per panel in `CHECK_LOG.txt`; any FAIL exits non-zero and stops the build.
- **Reading law:** a tier boundary is a descriptive CI-overlap fact under one judge, one epoch, one administration — pre-Gate-M, judges unvalidated, and (per the NIR caveat above) with intervals that may be too narrow on registered strata.

1.8 Rendering and verification

New quantities print at 3 decimals (ROUND_HALF_UP on the exact decimal); banked cells are copied at their banked 3-decimal renders. Figures carry the A-1 watermark and date stamp; seats appear in census order on every axis; no score-sorted axis exists. The markdown → HTML render is gated by `reports/mdhtml_cell_diff_20260828.py` (MATCH required) and the HTML → PDF pair by the digit-parity gate; the blind number-trace pass (NUMBER_TRACE) maps every printed value and load-bearing phrase to its source or recompute recipe.

---

§2 MFT — intuition-detection competence board (Family-C)

The MFT Family-C construct: within-school reasoning competence on Moral Foundations Theory's analytic architecture under adverse, complex, agentic, and multi-turn conditions. This chapter is one of three separate school readings; no seat's scores join across school readings (the only cross-chapter facts cited are school-neutral board-census counts, e.g. the FC-034 ceiling/floor census).

2.1 The banked board cells (COPIED — INV §C.2.1–.4, ADD §3)

Twenty-two seat-rows at school × setting grain: 21 rows of the 13-seat board (dual-judged substrates contribute one row per judge leg — **two readings of the same cells, never averaged**) plus the GCR1 own-grain row. Substrate blocks in sealed census order; scale 0–1; stratum n = 4 (S1p) / 16 (S2) / 14 (S3) unless noted.

seat (census order)	substrate	judge (leg)	mft S1p	mft S2	mft S3	notes

gpt-oss-20b	fc_gcr1_gptoss_20260828	claude-opus-5 (SINGLE-OOF)	0.500	0.622	0.118	
qwen3-235b-thinking	B3	claude-opus-5 (DUAL-OOF leg 1)	0.725	0.625	0.133	
deepseek-v4-flash	B3	claude-opus-5 (DUAL-OOF leg 1)	0.938	0.799	0.226	
mistral-small-3.2-24b	B3	claude-opus-5 (DUAL-OOF leg 1)	0.895	0.778	0.000	
qwen3-8b	B3	claude-opus-5 (DUAL-OOF leg 1)	0.480	0.738	0.000	
kimi-k2.6	B3	claude-opus-5 (DUAL-OOF leg 1)	0.730	0.889	0.476	
qwen3-235b-thinking	B3	gpt-5.6-sol (DUAL-OOF leg 2)	0.750	0.845	0.131	
deepseek-v4-flash	B3	gpt-5.6-sol (DUAL-OOF leg 2)	0.963	0.902	0.318	
mistral-small-3.2-24b	B3	gpt-5.6-sol (DUAL-OOF leg 2)	0.945	0.791	0.039	
qwen3-8b	B3	gpt-5.6-sol (DUAL-OOF leg 2)	0.730	0.914	0.000	
kimi-k2.6	B3	gpt-5.6-sol (DUAL-OOF leg 2)	0.750	0.981	0.477	
gemma-4-31b-it	B2	claude-opus-5 (DUAL-OOF leg 1)	0.955	0.846	0.177	
llama-4-scout	B2	claude-opus-5 (DUAL-OOF leg 1)	0.485	0.386	0.206	
glm-5.2	B2	claude-opus-5 (DUAL-OOF leg 1)	0.955	0.988	0.684	
gemma-4-31b-it	B2	gpt-5.6-sol (DUAL-OOF leg 2)	0.975	0.900	0.088	
llama-4-scout	B2	gpt-5.6-sol (DUAL-OOF leg 2)	0.930	0.813	0.366	
glm-5.2	B2	gpt-5.6-sol (DUAL-OOF leg 2)	0.980	0.978	0.621	
claude-opus-4-8	B1	gpt-5.6-sol (SINGLE-OOF)	1.000	0.966	0.682	
claude-sonnet-4-6	B1	gpt-5.6-sol (SINGLE-OOF)	0.980	0.972	0.822	
claude-haiku-4-5	B1	gpt-5.6-sol (SINGLE-OOF)	0.980	0.960	0.662	mft S3 0.662 = banked render of exact half 0.6625 (this recompute's ROUND_HALF_UP render is 0.663 — inside the banked envelope; banked value prints)
gpt-5.4	Frontier	claude-opus-5 (SINGLE-OOF)	0.955	0.463	0.516	
gpt-5.4-mini	Frontier	claude-opus-5 (SINGLE-OOF)	0.955	0.355	0.401	

Riders on this table. **GCR1 row:** substrate `fc_gcr1_gptoss_20260828`, post-A-14/A-16 epoch, single-OOF, subject custody **119/120 measured + 1 DECLARED** — rendered beside the 13-seat rows, never pooled with them. **B2 `claude-opus-5` rows:** the banked provenance tension FC-036 rides every citation — no published readout carries that per-seat surface; its primary source is the leg record tree via the verified derived-census recomputation (*"Carried, never resolved"*). **Stratum labels:** every S1p column is SMALL, every S2/S3 column MODERATE (pre-observation labels, not performance tiers).

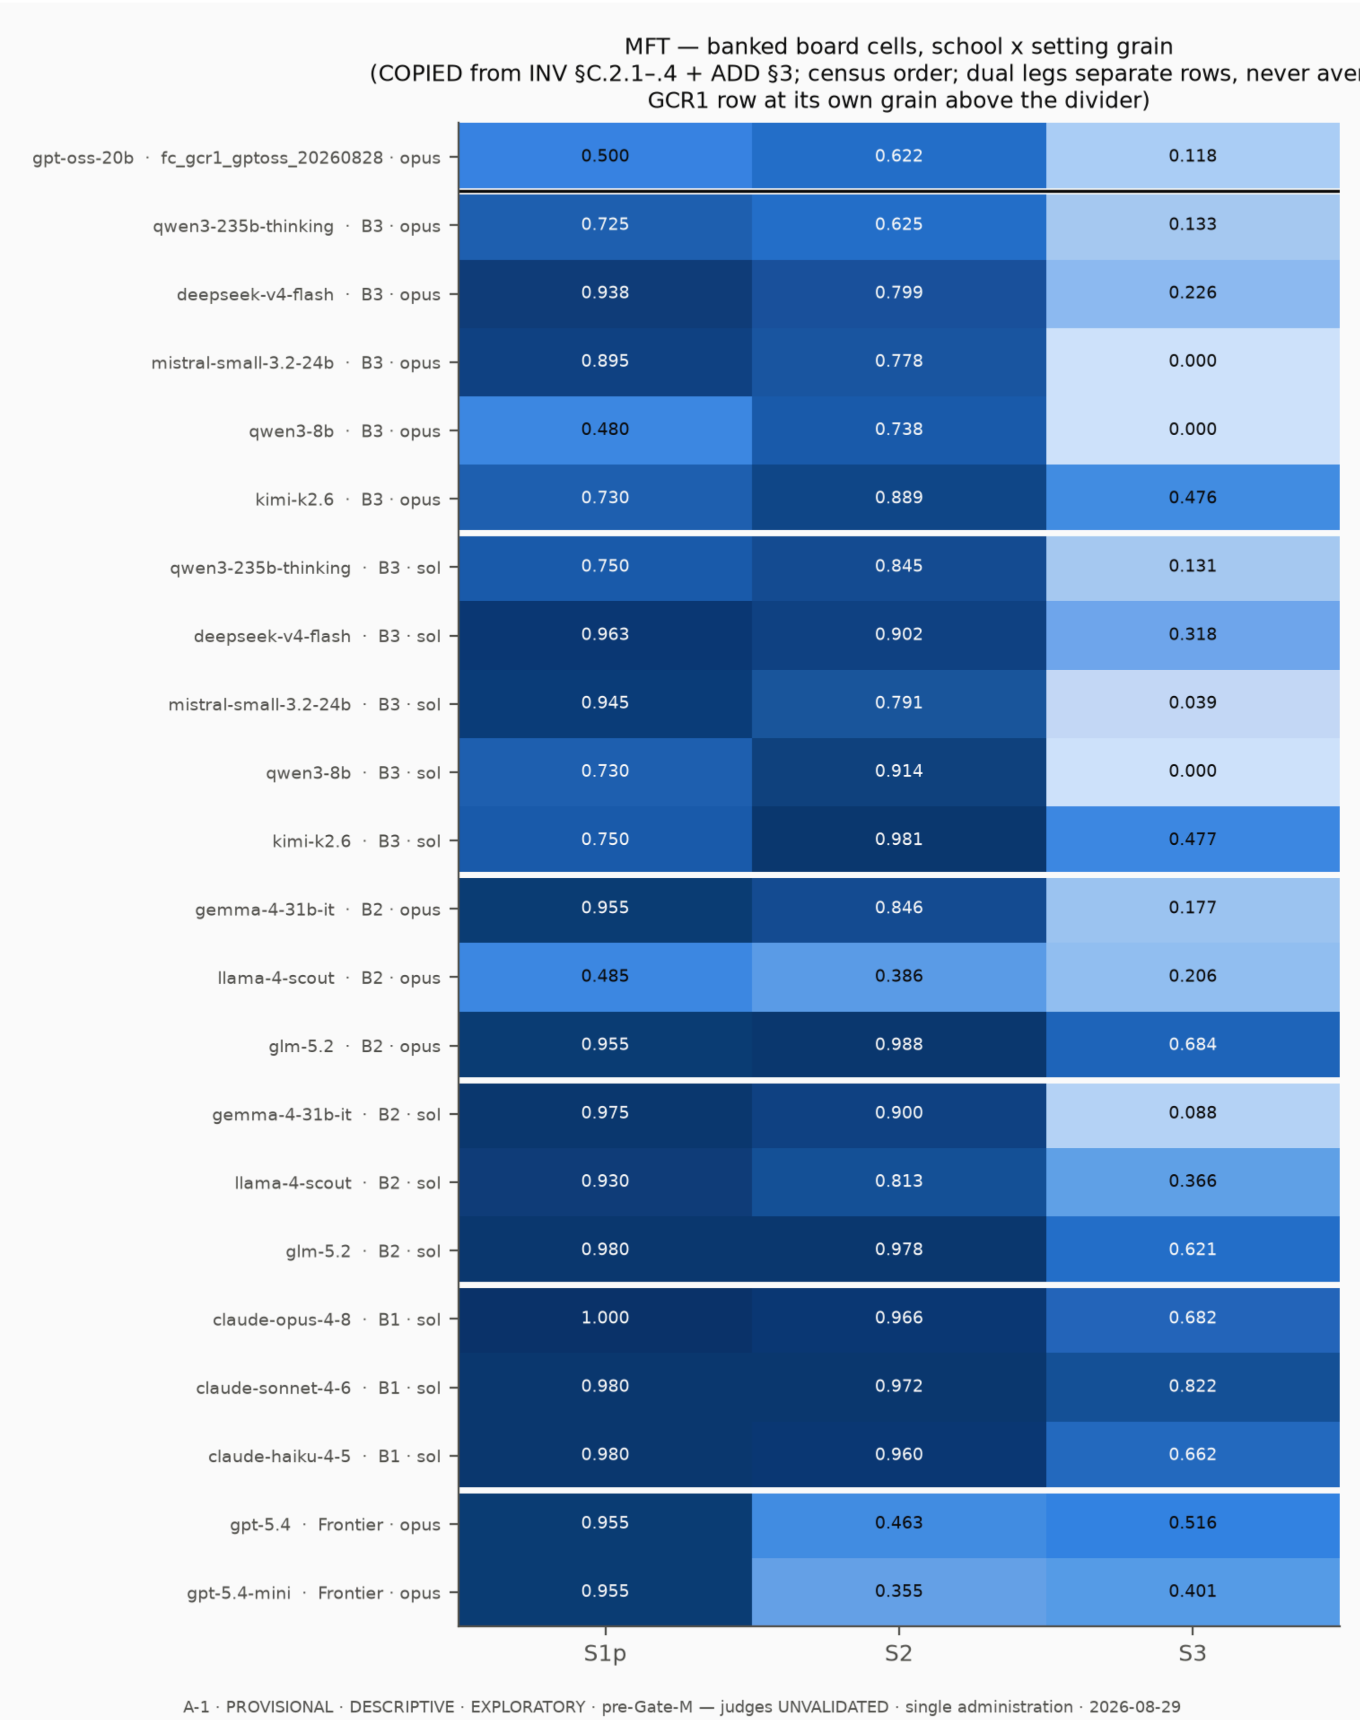
MFT — banked board cells, school x setting grain
(COPIED from INV §C.2.1–.4 + ADD §3; census order; dual legs separate rows, never aver
GCR1 row at its own grain above the divider)
gpt-oss-20b · fc_gcr1_gptoss_20260828 · opus
0.500
0.622
0.118
qwen3-235b-thinking · B3 · opus
0.725
0.625
0.133
deepseek-v4-flash · B3 · opus
0.938
0.799
0.226
mistral-small-3.2-24b · B3 · opus
0.895
0.778
0.000
qwen3-8b · B3 · opus
0.480
0.738
0.000
kimi-k2.6 · B3 · opus
0.730
0.889
0.476
qwen3-235b-thinking · B3 · sol
0.750
0.845
0.131
deepseek-v4-flash · B3 · sol
0.963
0.902
0.318
mistral-small-3.2-24b · B3 · sol
0.945
0.791
0.039
qwen3-8b · B3 · sol
0.730
0.914
0.000
kimi-k2.6 · B3 · sol
0.750
0.981
0.477
gemma-4-31b-it · B2 · opus
0.955
0.846
0.177
llama-4-scout · B2 · opus
0.485
0.386
0.206
glm-5.2 · B2 · opus
0.955
0.988
0.684
gemma-4-31b-it · B2 · sol
0.975
0.900
0.088
llama-4-scout · B2 · sol
0.930
0.813
0.366
glm-5.2 · B2 · sol
0.980
0.978
0.621
claude-opus-4-8 · B1 · sol
1.000
0.966
0.682
claude-sonnet-4-6 · B1 · sol
0.980
0.972
0.822
claude-haiku-4-5 · B1 · sol
0.980
0.960
0.662
gpt-5.4 · Frontier · opus
0.955
0.463
0.516
gpt-5.4-mini · Frontier · opus
0.955
0.355
0.401
S1p
S2
S3
A-1 · PROVISIONAL · DESCRIPTIVE · EXPLORATORY · pre-Gate-M — judges UNVALIDATED · single administration · 2026-08-29

2.2 Per-seat MFT means — `RECOMPUTED — NOT BANKED`

`RECOMPUTED — NOT BANKED` (D-263 §7F reach ruling). Recipe on every use: unweighted mean of the seat's three MFT setting-stratum means, per judge frame; 95% stratified item-bootstrap CI (§1.6). The composition includes the saturated, non-headline S1p stratum in equal weight. One number per seat **within MFT only** — never joined with any other school's reading.

frame	seat (census order)	school mean	95% CI	tier
F-A	`qwen3-235b-thinking`	0.575	[0.405, 0.716]	T1
F-A	`deepseek-v4-flash`	0.728	[0.642, 0.815]	T1
F-A	`mistral-small-3.2-24b`	0.592	[0.528, 0.648]	T1
F-A	`qwen3-8b`	0.548	[0.386, 0.645]	T1
F-A	`gemma-4-31b-it`	0.654	[0.596, 0.709]	T1
F-A	`llama-4-scout`	0.703	[0.621, 0.782]	T1
F-A	`glm-5.2`	0.860	[0.786, 0.928]	T1
F-A	`kimi-k2.6`	0.736	[0.561, 0.880]	T1
F-A	`claude-opus-4-8`	0.883	[0.802, 0.955]	T1
F-A	`claude-sonnet-4-6`	0.925	[0.893, 0.952]	T1
F-A	`claude-haiku-4-5`	0.868	[0.811, 0.914]	T1
F-B	`qwen3-235b-thinking`	0.494	[0.317, 0.642]	T1
F-B	`deepseek-v4-flash`	0.654	[0.571, 0.738]	T1
F-B	`mistral-small-3.2-24b`	0.558	[0.507, 0.601]	T1
F-B	`qwen3-8b`	0.406	[0.237, 0.573]	T1
F-B	`gemma-4-31b-it`	0.659	[0.590, 0.735]	T1
F-B	`llama-4-scout`	0.359	[0.182, 0.535]	T1
F-B	`glm-5.2`	0.875	[0.803, 0.939]	T1
F-B	`kimi-k2.6`	0.698	[0.517, 0.847]	T1
F-C	`gpt-5.4`	0.645	[0.545, 0.743]	T1
F-C	`gpt-5.4-mini`	0.570	[0.474, 0.670]	T1
F-D	`gpt-oss-20b`	0.413	[0.232, 0.592]	own-grain

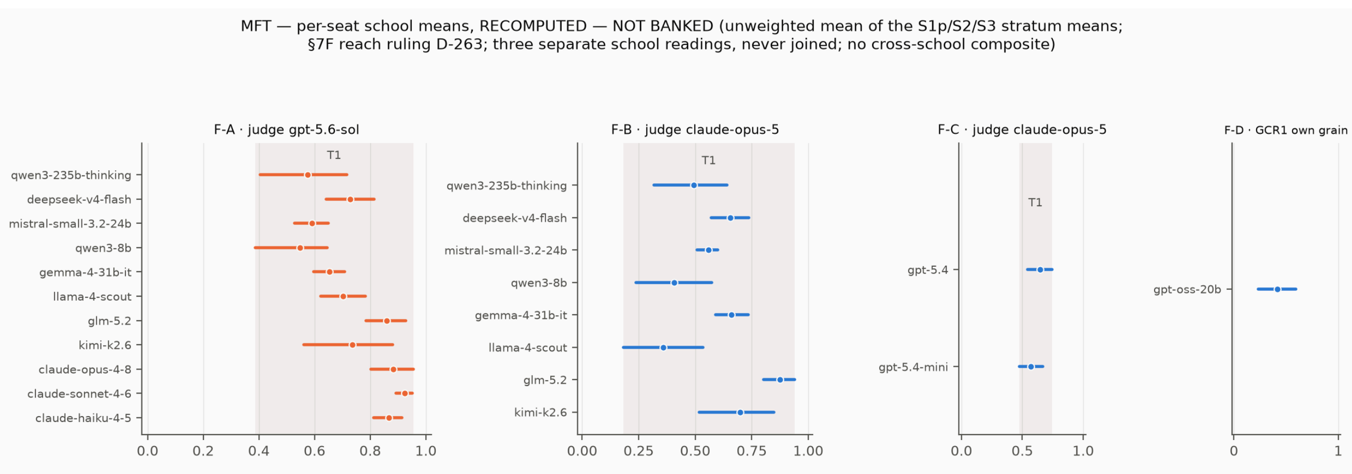

2.3 CI-backed tiers per same-judge frame

Under the frozen rule (§1.7), the MFT panels resolve as follows:

frame	setting	tier	members (census order; no within-tier order)
F-A	S2	T1	`qwen3-235b-thinking`, `deepseek-v4-flash`, `mistral-small-3.2-24b`, `qwen3-8b`, `gemma-4-31b-it`, `llama-4-scout`, `glm-5.2`, `kimi-`

			k2.6, claude-opus-4-8, claude-sonnet-4-6, claude-haiku-4-5
F-A	S3	T1	qwen3-235b-thinking, deepseek-v4-flash, mistral-small-3.2-24b, qwen3-8b, gemma-4-31b-it, llama-4-scout, glm-5.2, kimi-k2.6, claude-opus-4-8, claude-sonnet-4-6, claude-haiku-4-5
F-A	S1	T1	qwen3-235b-thinking, deepseek-v4-flash, mistral-small-3.2-24b, qwen3-8b, gemma-4-31b-it, llama-4-scout, glm-5.2, kimi-k2.6, claude-opus-4-8, claude-sonnet-4-6, claude-haiku-4-5
F-B	S2	T1	qwen3-235b-thinking, deepseek-v4-flash, mistral-small-3.2-24b, qwen3-8b, gemma-4-31b-it, llama-4-scout, glm-5.2, kimi-k2.6
F-B	S3	T1	qwen3-235b-thinking, deepseek-v4-flash, mistral-small-3.2-24b, qwen3-8b, gemma-4-31b-it, llama-4-scout, glm-5.2, kimi-k2.6
F-B	S1	T1	qwen3-235b-thinking, deepseek-v4-flash, mistral-small-3.2-24b, qwen3-8b, gemma-4-31b-it, llama-4-scout, glm-5.2, kimi-k2.6
F-C	S2	T1	gpt-5.4, gpt-5.4-mini
F-C	S3	T1	gpt-5.4, gpt-5.4-mini
F-C	S1	T1	gpt-5.4, gpt-5.4-mini

Every MFT panel — every setting in F-A, F-B, and F-C, and every school-mean panel — yields a single tier, T1. The same holds in §3 and §4: at the current provisional 14-seat state, no school × setting × frame panel CI-separates into more than one tier under the declared CI-overlap chaining rule. This is the report's central comparative result, and it is a **chained band, not an equivalence statement**: the frozen rule's transitive closure merges seats whose own intervals do not overlap whenever intermediate seats bridge them, and with 8–11 seats spread quasi-continuously over the score range, bridges exist in every panel. Floor- and ceiling-anchored cells extend the chains further (a bootstrap interval on a mostly-zero stratum reaches 0.000; a saturated stratum sits degenerate at 1.000).

Three riders keep this result honestly scoped. First, the CIs are **anti-conservative** wherever a stratum contains registered non-independent items (NIR R1/R8 record a declared template class on the mft S2 slice, among others), so the true uncertainty is wider and the single-tier outcome is, if anything, understated. Second, the S1p panels carry the SMALL-stratum caveat: n = 4, coarse bootstrap support, no headline. Third, this is a **descriptive pre-Gate-M reading under one unvalidated judge per frame**; it neither executes nor bears on the sealed tier-separability hypothesis, which is a confirmatory Gate-M-gated question on a different analysis path.

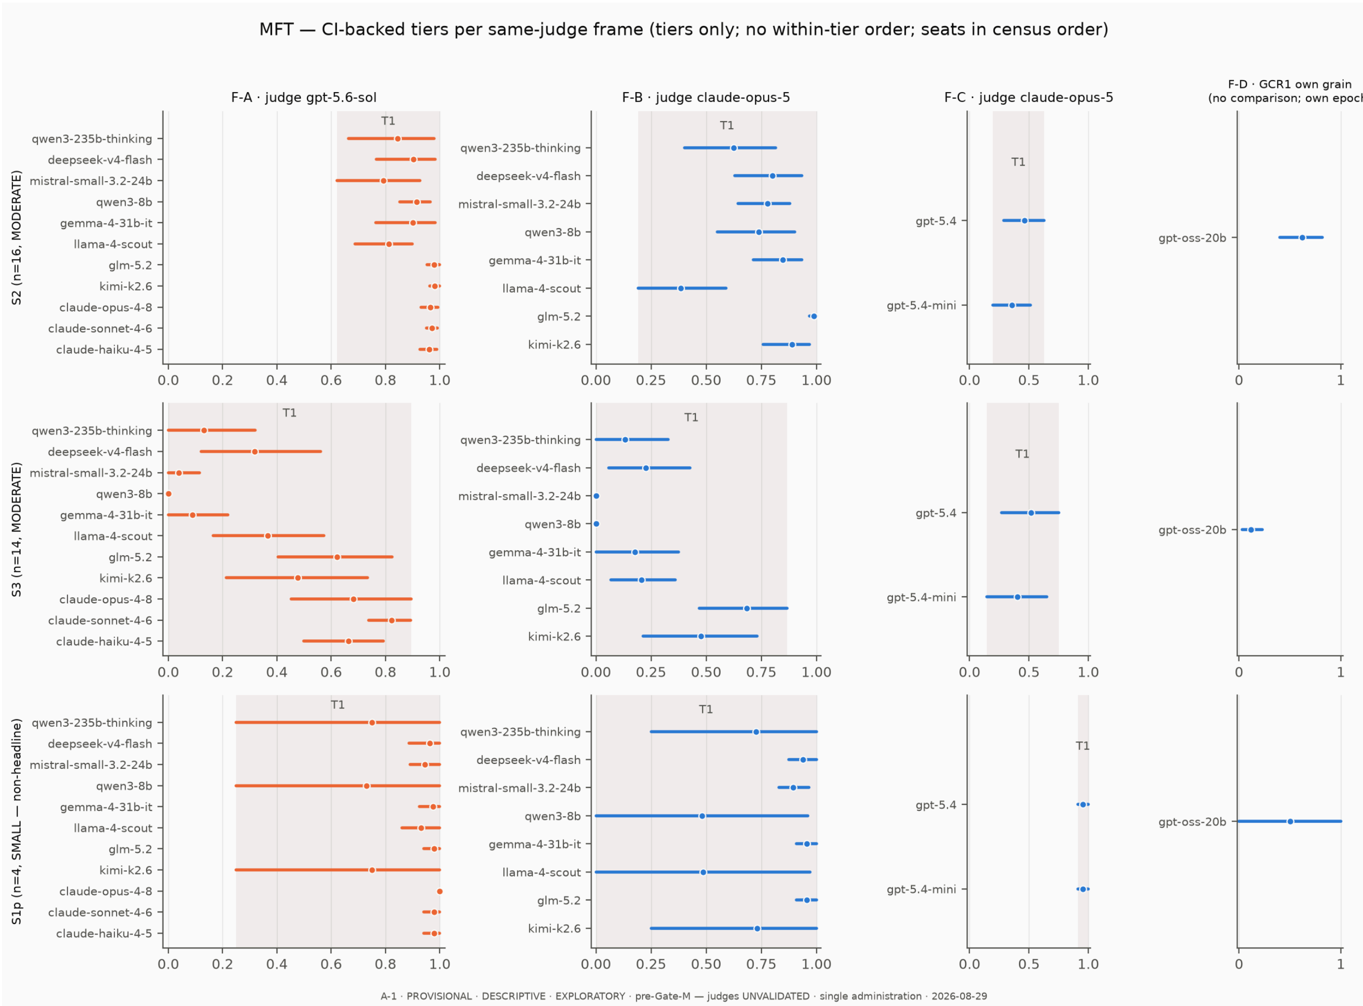

2.4 What the MFT board shows (description at tier grain)

S2 (agentic, n = 16) sits high in every frame. In F-A the eleven seats' S2 cells span 0.791–0.981; in F-B the eight seats span 0.386–0.988. The two frames render side by side only: they differ by judge, no cross-frame comparison or difference is made, and any such number would be judge-confounded. In F-C the two OpenAI-substrate seats' S2 cells (0.463, 0.355) sit in the middle of the scale, under that substrate's own inverted setting order (a banked shape statement: S2 lowest within that leg — never a ranking).

S3 (multi-turn adversity, n = 14) is where MFT competence drops. Three of the board's zero cells are MFT S3 cells (`mistral-small-3.2-24b` under F-B; `qwen3-8b` under both F-A and F-B — FLOOR position labels of record), and every frame's MFT S3 column sits well below its S2 column. The S3 band is wide in both main frames (F-A: 0.000–0.822; F-B: 0.000–0.684), which is precisely where the chaining rule absorbs would-be tier boundaries: the panel is a continuum, not two clumps.

S1p (n = 4, SMALL) is near-saturated at the top of the roster — eight of eleven F-A seats sit at or above 0.930, with `claude-opus-4-8` at the board's only non-schwartz ceiling cell (1.000) — and correspondingly uninformative for separation; its tier panels are rendered for completeness only.

The GCR1 seat, read beside (never inside) the frames: `gpt-oss-20b`'s MFT cells are 0.500 (S1p) / 0.622 (S2) / 0.118 (S3), MFT school mean 0.413 [0.232, 0.592] (`RECOMPUTED — NOT BANKED`, recipe §1.6) — at its own substrate grain, own epoch, single-OOF `claude-opus-5` judging. Its S3 value continues the board-wide MFT-S3-adversity pattern; no comparative statement attaches to it.

Tier-grain summary. Within any one frame, the licensed comparison yields: no MFT CI-separation between any seats at any setting — the 14 measured seats' MFT readings form one chained band per panel, with the spread (floors at S3, saturation at S1p, high S2) carried by the descriptive shape above, not by tier boundaries.

§3 Schwartz — value-architecture competence board (Family-C)

The Schwartz Family-C construct: within-school reasoning competence on the refined Schwartz value-priority architecture under adverse, complex, agentic, and multi-turn conditions. One of three separate school readings; no seat's scores join across school readings (cross-chapter citations are school-neutral board-census counts only). This is the competence board — no Family-P value-priority profile, and no per-seat 19-value vector, appears here or anywhere (G2; the priority-vector key is never opened).

3.1 The banked board cells (COPIED — INV §C.2.1–.4, ADD §3)

seat (census order)	substrate	judge (leg)	schwartz S1p	schwartz S2	schwartz S3	notes
`gpt-oss-20b`	fc_gcr1_gptoss_20260828	`claude-opus-5` (SINGLE-OOF)	0.750	0.844	0.412	sch S3 over n=13 of 14 (declared cell, never imputed)
`qwen3-235b-thinking`	B3	`claude-opus-5` (DUAL-OOF leg 1)	0.719	0.709	0.089	
`deepseek-v4-flash`	B3	`claude-opus-5` (DUAL-OOF leg 1)	1.000	0.947	0.329	
`mistral-small-3.2-24b`	B3	`claude-opus-5` (DUAL-OOF leg 1)	0.725	0.759	0.346	
`qwen3-8b`	B3	`claude-opus-5` (DUAL-OOF leg 1)	0.500	0.673	0.257	
`kimi-k2.6`	B3	`claude-opus-5` (DUAL-OOF leg 1)	0.750	0.944	0.473	
`qwen3-235b-thinking`	B3	`gpt-5.6-sol` (DUAL-OOF leg 2)	0.500	0.789	0.093	
`deepseek-v4-flash`	B3	`gpt-5.6-sol` (DUAL-OOF leg 2)	1.000	0.933	0.491	
`mistral-small-3.2-24b`	B3	`gpt-5.6-sol` (DUAL-OOF leg 2)	0.750	0.833	0.468	
`qwen3-8b`	B3	`gpt-5.6-sol` (DUAL-OOF leg 2)	0.750	0.842	0.321	
`kimi-k2.6`	B3	`gpt-5.6-sol` (DUAL-OOF leg 2)	1.000	0.936	0.379	
`gemma-4-31b-it`	B2	`claude-opus-5` (DUAL-OOF leg 1)	1.000	0.861	0.516	
`llama-4-scout`	B2	`claude-opus-5` (DUAL-OOF leg 1)	0.919	0.867	0.464	
`glm-5.2`	B2	`claude-opus-5` (DUAL-OOF leg 1)	0.750	0.964	0.686	

`gemma-4-31b-it`	B2	`gpt-5.6-sol` (DUAL-OOF leg 2)	1.000	0.880	0.500	
`llama-4-scout`	B2	`gpt-5.6-sol` (DUAL-OOF leg 2)	1.000	0.902	0.563	sch S3 0.563 of record (CONFLICT-A / TENSION-1: exact mean 0.5625)
`glm-5.2`	B2	`gpt-5.6-sol` (DUAL-OOF leg 2)	1.000	0.923	0.871	
`claude-opus-4-8`	B1	`gpt-5.6-sol` (SINGLE-OOF)	1.000	0.817	0.386	
`claude-sonnet-4-6`	B1	`gpt-5.6-sol` (SINGLE-OOF)	0.500	0.938	0.891	
`claude-haiku-4-5`	B1	`gpt-5.6-sol` (SINGLE-OOF)	0.750	0.909	0.643	
`gpt-5.4`	Frontier	`claude-opus-5` (SINGLE-OOF)	1.000	0.565	0.582	sch S2 over n=15 (disclosed no-verdict, never imputed)
`gpt-5.4-mini`	Frontier	`claude-opus-5` (SINGLE-OOF)	1.000	0.466	0.505	

Riders on this table. **The two declared no-verdicts both sit in this school:** the lane's only board no-verdict, `sch-s2-003` × `gpt-5.4` (judge safeguard refusal — that stratum is n = 15, disclosed, never imputed), and the GCR1 declared cell `sch-s3-013` (subject-side `empty_output` — n = 13 of 14). **GCR1 row:** own substrate, post-A-14/A-16 epoch, single-OOF, 119/120 + 1-declared custody rider. **B2 `claude-opus-5` rows:** the FC-036 provenance tension rides every citation (no published readout carries that per-seat surface; primary source is the leg record tree via the verified derived-census recomputation — carried, never resolved). **CONFLICT-A** rides `llama-4-scout`'s S3 cell under the `gpt-5.6-sol` leg: 0.563 is the edition-of-record render of the exact mean 0.5625 (three readout renders print .562 — the banked TENSION-1 rounding divergence; never shown as equal, never averaged).

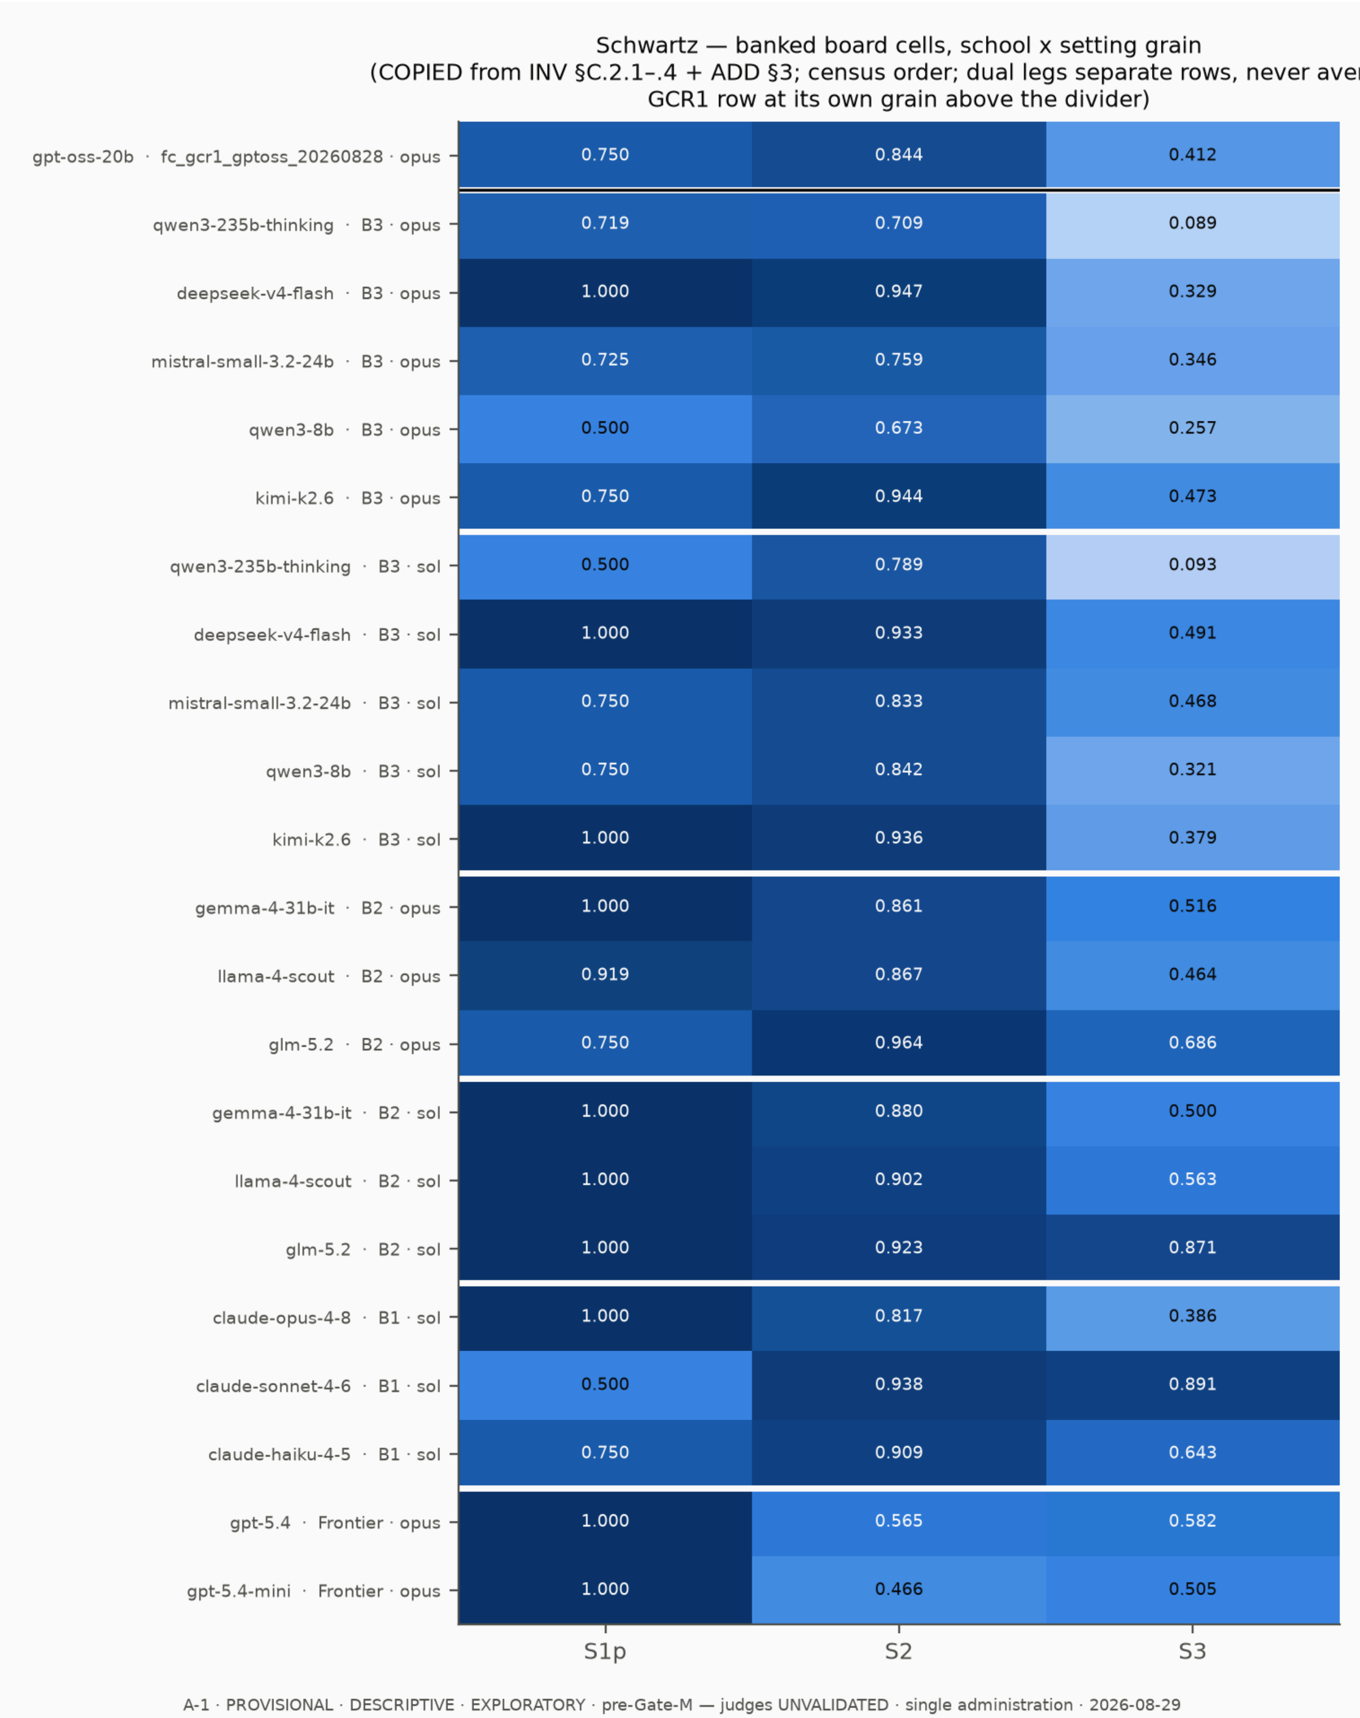

3.2 Per-seat Schwartz means — `RECOMPUTED — NOT BANKED`

`RECOMPUTED — NOT BANKED` (D-263 §7F reach ruling; recipe §1.6, composition disclosed — the near-saturated S1p stratum contributes 1/3 weight). One number per seat **within Schwartz only**.

frame	seat (census order)	school mean	95% CI	tier
F-A	`qwen3-235b-thinking`	0.461	[0.284, 0.640]	T1
F-A	`deepseek-v4-flash`	0.808	[0.731, 0.883]	T1
F-A	`mistral-small-3.2-24b`	0.684	[0.508, 0.824]	T1
F-A	`qwen3-8b`	0.638	[0.459, 0.783]	T1
F-A	`gemma-4-31b-it`	0.793	[0.712, 0.868]	T1
F-A	`llama-4-scout`	0.821	[0.750, 0.887]	T1
F-A	`glm-5.2`	0.932	[0.880, 0.966]	T1
F-A	`kimi-k2.6`	0.772	[0.695, 0.851]	T1
F-A	`claude-opus-4-8`	0.734	[0.641, 0.828]	T1
F-A	`claude-sonnet-4-6`	0.776	[0.607, 0.946]	T1
F-A	`claude-haiku-4-5`	0.767	[0.591, 0.903]	T1
F-B	`qwen3-235b-thinking`	0.506	[0.338, 0.638]	T1
F-B	`deepseek-v4-flash`	0.758	[0.684, 0.839]	T1
F-B	`mistral-small-3.2-24b`	0.610	[0.438, 0.749]	T1
F-B	`qwen3-8b`	0.477	[0.294, 0.663]	T1
F-B	`gemma-4-31b-it`	0.792	[0.708, 0.871]	T1
F-B	`llama-4-scout`	0.750	[0.682, 0.817]	T1
F-B	`glm-5.2`	0.800	[0.624, 0.938]	T1
F-B	`kimi-k2.6`	0.722	[0.545, 0.861]	T1
F-C	`gpt-5.4`	0.716	[0.626, 0.799]	T1
F-C	`gpt-5.4-mini`	0.657	[0.566, 0.747]	T1
F-D	`gpt-oss-20b`	0.668	[0.495, 0.803]	own-grain

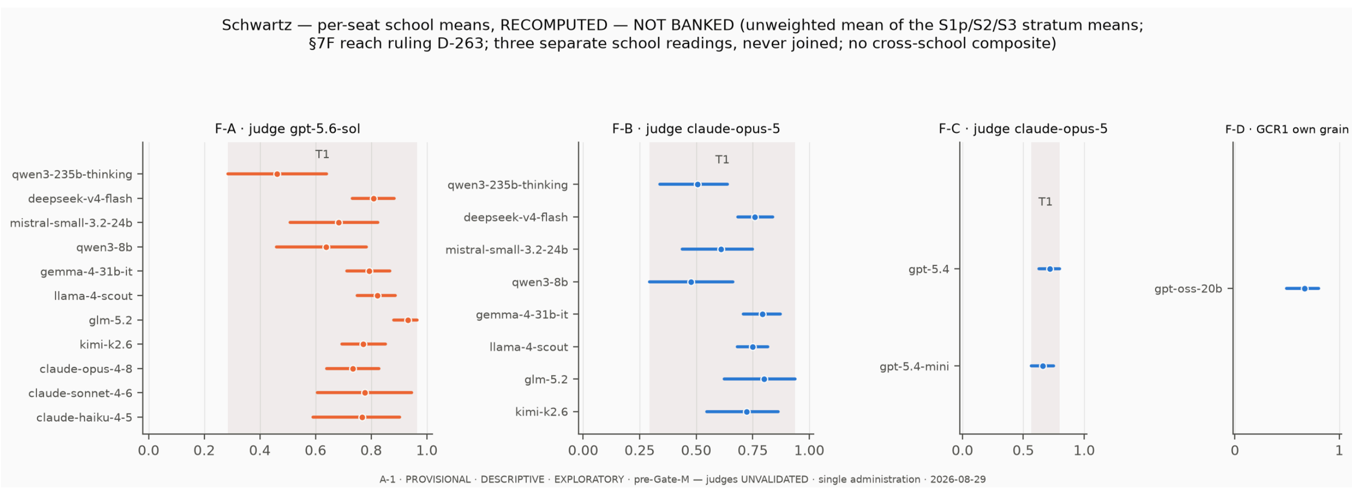

3.3 CI-backed tiers per same-judge frame

frame	setting	tier	members (census order; no within-tier order)
F-A	S2	T1	`qwen3-235b-thinking`, `deepseek-v4-flash`, `mistral-small-3.2-24b`, `qwen3-8b`, `gemma-4-31b-it`, `llama-4-scout`, `glm-5.2`, `kimi-k2.6`, `claude-opus-4-8`, `claude-sonnet-4-6`, `claude-haiku-4-5`
F-A	S3	T1	`qwen3-235b-thinking`, `deepseek-v4-flash`, `mistral-small-3.2-24b`, `qwen3-8b`, `gemma-4-31b-it`, `llama-4-scout`, `glm-5.2`, `kimi-k2.6`, `claude-opus-4-8`, `claude-sonnet-4-6`, `claude-haiku-4-5`
F-A	S1	T1	`qwen3-235b-thinking`, `deepseek-v4-flash`, `mistral-small-3.2-24b`, `qwen3-8b`, `gemma-4-31b-it`, `llama-4-scout`, `glm-5.2`, `kimi-`

			k2.6, claude-opus-4-8, claude-sonnet-4-6, claude-haiku-4-5
F-B	S2	T1	qwen3-235b-thinking, deepseek-v4-flash, mistral-small-3.2-24b, qwen3-8b, gemma-4-31b-it, llama-4-scout, glm-5.2, kimi-k2.6
F-B	S3	T1	qwen3-235b-thinking, deepseek-v4-flash, mistral-small-3.2-24b, qwen3-8b, gemma-4-31b-it, llama-4-scout, glm-5.2, kimi-k2.6
F-B	S1	T1	qwen3-235b-thinking, deepseek-v4-flash, mistral-small-3.2-24b, qwen3-8b, gemma-4-31b-it, llama-4-scout, glm-5.2, kimi-k2.6
F-C	S2	T1	gpt-5.4, gpt-5.4-mini
F-C	S3	T1	gpt-5.4, gpt-5.4-mini
F-C	S1	T1	gpt-5.4, gpt-5.4-mini

As in §2.3, **every Schwartz panel yields the single tier T1** under the frozen rule — a chained band, not an equivalence statement. The Schwartz chains are anchored at the top by this school's distinctive saturation: **ten of the board's eleven exact-1.000 cells are schwartz S1p cells** (the ratified count of record — E-1, RULED at D-178: PROPAGATE; the superseded "nine" reading is disclosed, not silently dropped). A saturated stratum sits as a degenerate zero-width interval at 1.000, and the S1p panels' chains form beneath it. The anti-conservativeness rider (NIR class-grain schwartz domain-echo sets, R1/R8) and the SMALL-stratum caveat apply as in §2.3.

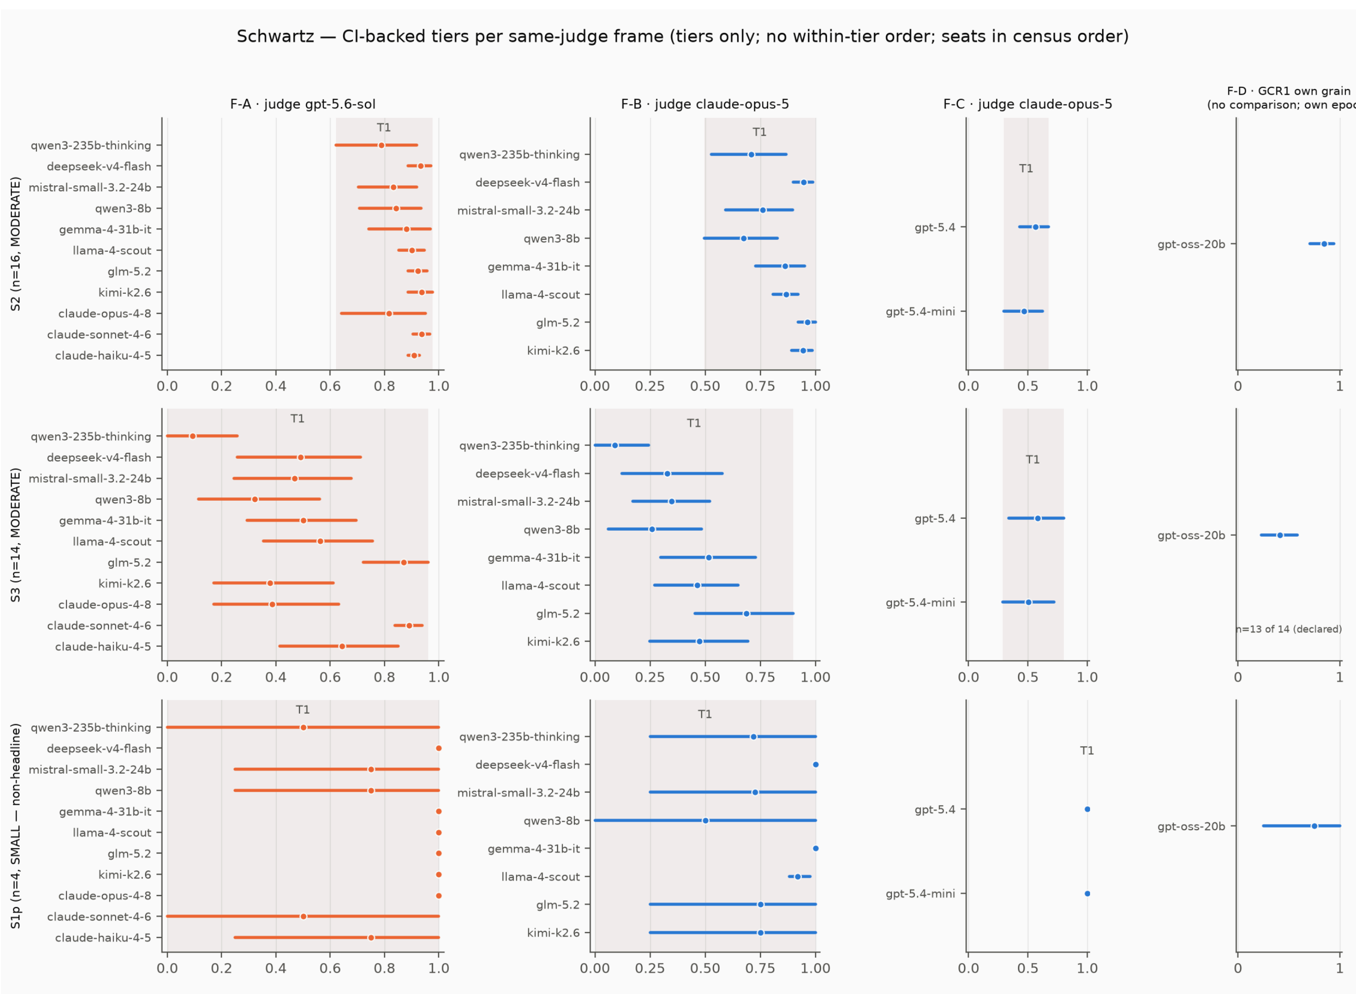

3.4 What the Schwartz board shows (description at tier grain)

S1p (n = 4, SMALL) is the saturation stratum. Six of the eleven F-A seats, two of eight in F-B, and both F-C seats sit at exactly 1.000 — floor-and-ceiling strata are *"uninformative for discrimination and are disclosed as design inputs, never as seat verdicts."* The stratum also carries the board's widest SMALL-stratum artifacts (e.g. `claude-sonnet-4-6` at 0.500, where a single cell moves the mean visibly); no headline is drawn from S1.

S2 (n = 16) is this school's strongest column. F-A spans 0.789–0.938; F-B spans 0.673–0.964; the two F-C cells are 0.565 (n = 15) and 0.466, under that substrate's banked inverted setting order (shape statement, not a ranking).

S3 (n = 14) drops and widens. F-A spans 0.093–0.891 and F-B spans 0.089–0.686 — in both frames the low end is `qwen3-235b-thinking`'s S3 cell and the band up from it is quasi-continuous, which is exactly the chaining geometry: pairwise-disjoint extremes bridged by intermediate seats. (Naming the extremes' values describes banked cells; it asserts no within-tier order and no pairwise verdict.)

The GCR1 seat, read beside the frames: `gpt-oss-20b`'s Schwartz cells are 0.750 (S1p) / 0.844 (S2) / 0.412 (S3, n = 13 of 14 — the declared cell), Schwartz school mean 0.668 [0.495, 0.803] (`RECOMPUTED — NOT BANKED`, recipe §1.6), own grain, own epoch. Its shape follows the school-wide pattern (S2 strongest, S3 weakest); no comparative statement attaches.

Tier-grain summary. Within any one frame, no Schwartz CI-separation between any seats at any setting: one chained band per panel, top-anchored by S1p saturation, with the S3 column carrying the school's spread.

§4 Kohlberg — reasoning-structure competence board (Family-C)

The Kohlberg Family-C construct: within-school competence on the cognitive-developmental tradition's analytic architecture — reasoning-structure scoring under adverse, complex, agentic, and multi-turn conditions. One of three separate school readings; no seat's scores join across school readings (cross-chapter citations are school-neutral board-census counts and instrument-level column shapes only, never seat scores). No developmental, stage, level, or maturity attribution is made about any seat anywhere in this chapter.

4.1 The banked board cells (COPIED — INV §C.2.1–.4, ADD §3)

seat (census order)	substrate	judge (leg)	kohlberg S1p	kohlberg S2	kohlberg S3	notes
`gpt-oss-20b`	fc_gcr1_gptoss_20260828	`claude-opus-5` (SINGLE-OOF)	0.808	0.866	0.498	koh S1p 0.808 = ROUND_HALF_UP render of exact half 0.8075
`qwen3-235b-thinking`	B3	`claude-opus-5` (DUAL-OOF leg 1)	0.613	0.684	0.425	
`deepseek-v4-flash`	B3	`claude-opus-5` (DUAL-OOF leg 1)	0.184	0.933	0.731	
`mistral-small-3.2-24b`	B3	`claude-opus-5` (DUAL-OOF leg 1)	0.223	0.643	0.225	
`qwen3-8b`	B3	`claude-opus-5` (DUAL-OOF leg 1)	0.000	0.688	0.151	

kimi-k2.6	B3	claude-opus-5 (DUAL-OOF leg 1)	0.640	0.930	0.806	
qwen3-235b-thinking	B3	gpt-5.6-sol (DUAL-OOF leg 2)	0.780	0.653	0.531	
deepseek-v4-flash	B3	gpt-5.6-sol (DUAL-OOF leg 2)	0.835	0.885	0.657	
mistral-small-3.2-24b	B3	gpt-5.6-sol (DUAL-OOF leg 2)	0.703	0.566	0.294	
qwen3-8b	B3	gpt-5.6-sol (DUAL-OOF leg 2)	0.585	0.703	0.331	
kimi-k2.6	B3	gpt-5.6-sol (DUAL-OOF leg 2)	0.835	0.895	0.701	
gemma-4-31b-it	B2	claude-opus-5 (DUAL-OOF leg 1)	0.390	0.857	0.494	
llama-4-scout	B2	claude-opus-5 (DUAL-OOF leg 1)	0.585	0.673	0.391	
glm-5.2	B2	claude-opus-5 (DUAL-OOF leg 1)	0.835	0.963	0.972	
gemma-4-31b-it	B2	gpt-5.6-sol (DUAL-OOF leg 2)	0.725	0.881	0.656	
llama-4-scout	B2	gpt-5.6-sol (DUAL-OOF leg 2)	0.725	0.413	0.429	
glm-5.2	B2	gpt-5.6-sol (DUAL-OOF leg 2)	0.863	0.867	0.826	
claude-opus-4-8	B1	gpt-5.6-sol (SINGLE-OOF)	0.890	0.979	0.874	
claude-sonnet-4-6	B1	gpt-5.6-sol (SINGLE-OOF)	0.890	0.920	0.873	
claude-haiku-4-5	B1	gpt-5.6-sol (SINGLE-OOF)	0.672	0.804	0.808	koh S1p 0.672 = banked render of exact half 0.6725 (this recompute's ROUND_HALF_UP render is 0.673 — inside the banked envelope; banked value prints)
gpt-5.4	Frontier	claude-opus-5 (SINGLE-OOF)	0.918	0.434	0.900	koh S1p 0.918 of record (CONFLICT-B: exact mean 0.9175, ADJ-2)
gpt-5.4-mini	Frontier	claude-opus-5 (SINGLE-OOF)	0.863	0.295	0.838	

Riders on this table. **CONFLICT-B** rides `gpt-5.4`'s S1p cell: 0.918 is the edition-of-record render of the exact mean 0.9175 (three readouts print .917); per ADJ-2 this report prints the edition-of-record value only. `claude-haiku-4-5`'s S1p 0.672 is the banked render of the exact half 0.6725 (see the §6 render-divergence disclosure). The GCR1 row's S1p 0.808 is the ROUND_HALF_UP render of the exact half 0.8075 (ADD §3, the board-of-record convention). **GCR1 row:** own substrate, post-A-14/A-16 epoch, single-OOF, 119/120 + 1-declared custody rider. **B2 `claude-opus-5` rows:** the FC-036 provenance tension rides every citation, carried, never resolved. Two B3 seats carry re-administered kohlberg cells (the FC-057 remediation — four cells, three records quarantined, re-administered fresh; a missing record is a pending cell, never a zero); the values printed are the re-administered ones.

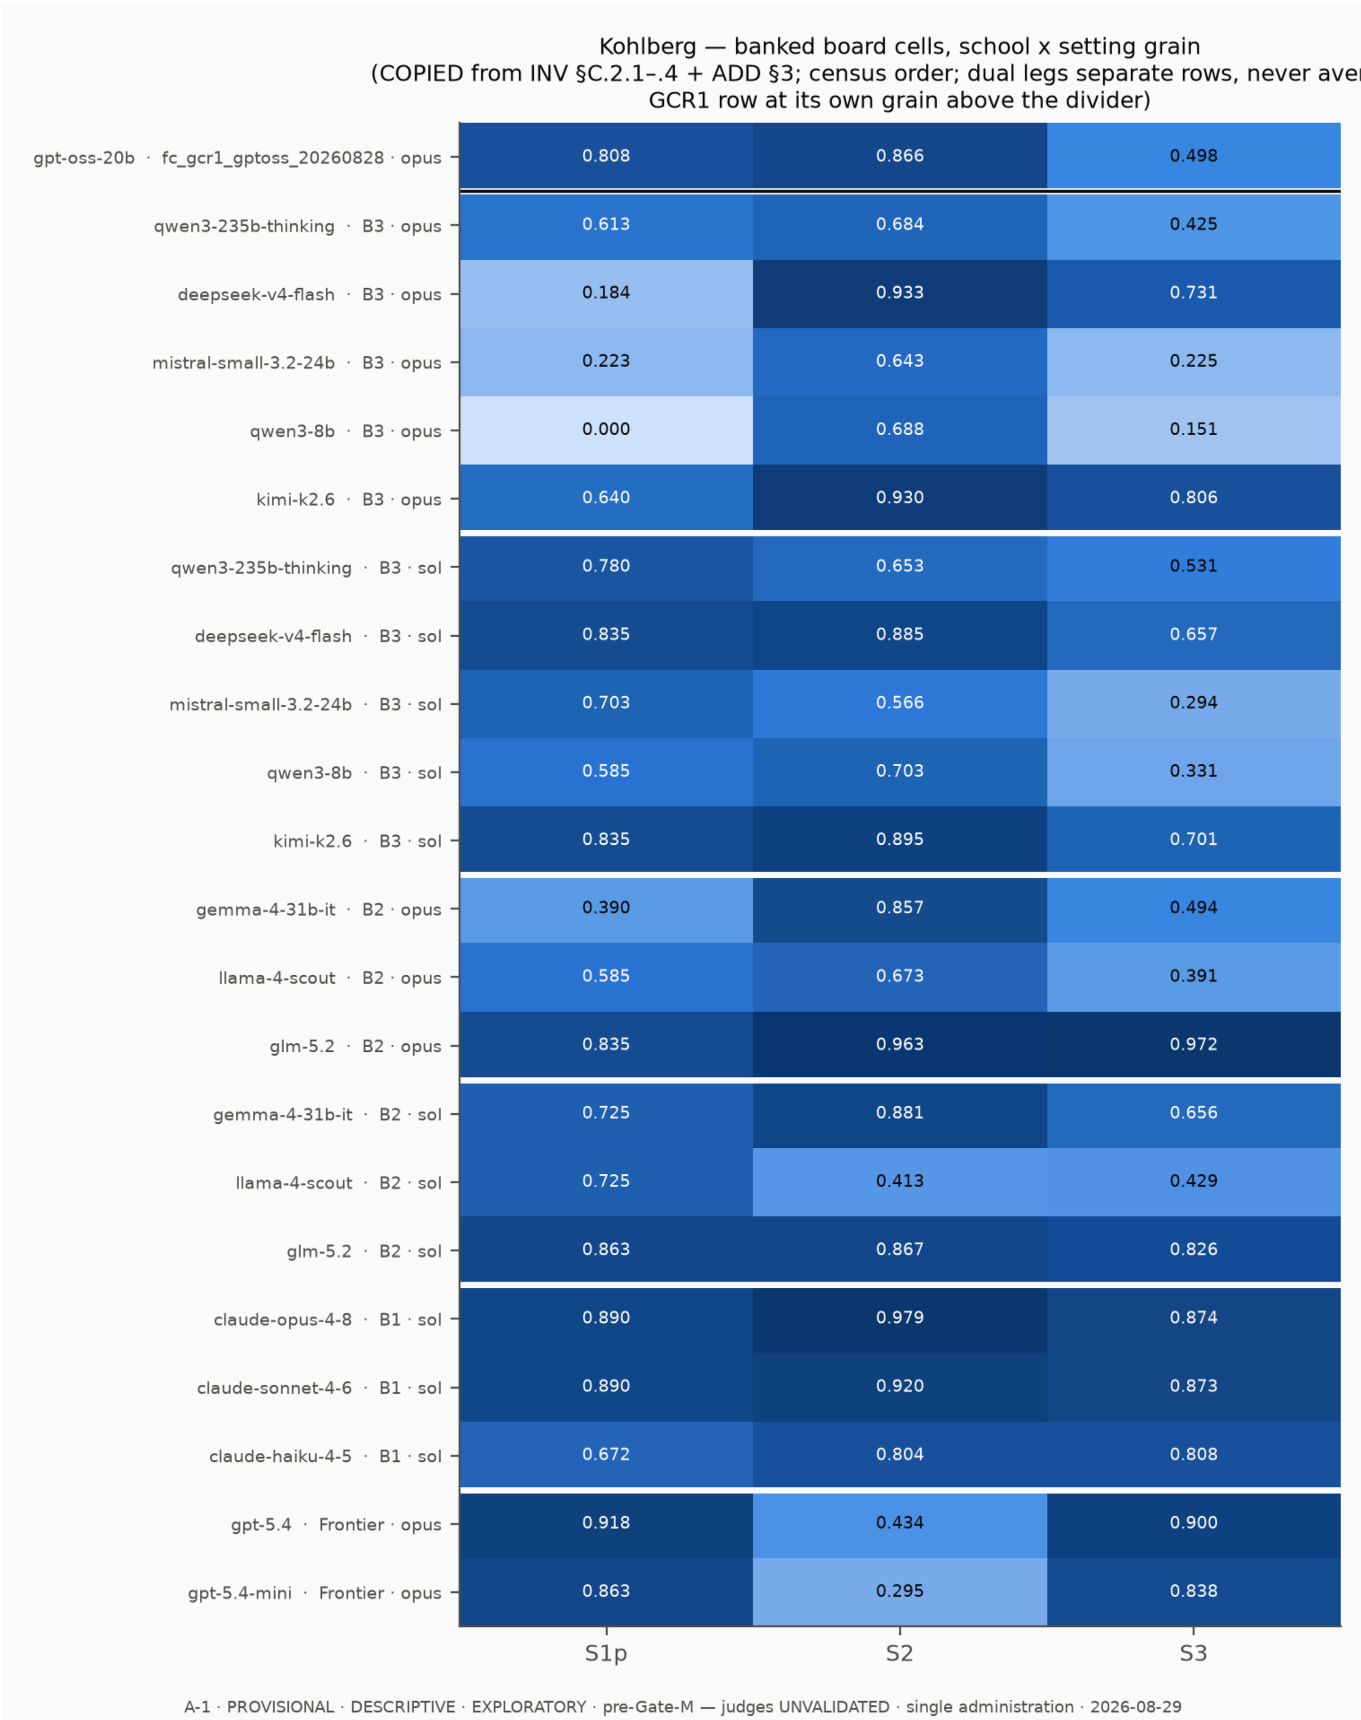

4.2 Per-seat Kohlberg means — `RECOMPUTED — NOT BANKED`

`RECOMPUTED — NOT BANKED` (D-263 §7F reach ruling; recipe §1.6). One number per seat **within Kohlberg only**.

frame	seat (census order)	school mean	95% CI	tier
F-A	`qwen3-235b-thinking`	0.655	[0.548, 0.759]	T1
F-A	`deepseek-v4-flash`	0.792	[0.700, 0.873]	T1
F-A	`mistral-small-3.2-24b`	0.521	[0.417, 0.627]	T1
F-A	`qwen3-8b`	0.540	[0.383, 0.673]	T1
F-A	`gemma-4-31b-it`	0.754	[0.664, 0.834]	T1
F-A	`llama-4-scout`	0.523	[0.413, 0.633]	T1
F-A	`glm-5.2`	0.852	[0.759, 0.931]	T1
F-A	`kimi-k2.6`	0.810	[0.717, 0.893]	T1
F-A	`claude-opus-4-8`	0.915	[0.857, 0.959]	T1
F-A	`claude-sonnet-4-6`	0.894	[0.821, 0.954]	T1
F-A	`claude-haiku-4-5`	0.762	[0.604, 0.890]	T1
F-B	`qwen3-235b-thinking`	0.574	[0.413, 0.714]	T1
F-B	`deepseek-v4-flash`	0.616	[0.502, 0.750]	T1
F-B	`mistral-small-3.2-24b`	0.364	[0.231, 0.524]	T1
F-B	`qwen3-8b`	0.280	[0.203, 0.356]	T1
F-B	`gemma-4-31b-it`	0.580	[0.427, 0.733]	T1
F-B	`llama-4-scout`	0.550	[0.402, 0.675]	T1
F-B	`glm-5.2`	0.924	[0.900, 0.946]	T1
F-B	`kimi-k2.6`	0.792	[0.643, 0.906]	T1
F-C	`gpt-5.4`	0.750	[0.668, 0.826]	T1
F-C	`gpt-5.4-mini`	0.665	[0.587, 0.743]	T1
F-D	`gpt-oss-20b`	0.724	[0.659, 0.785]	own-grain

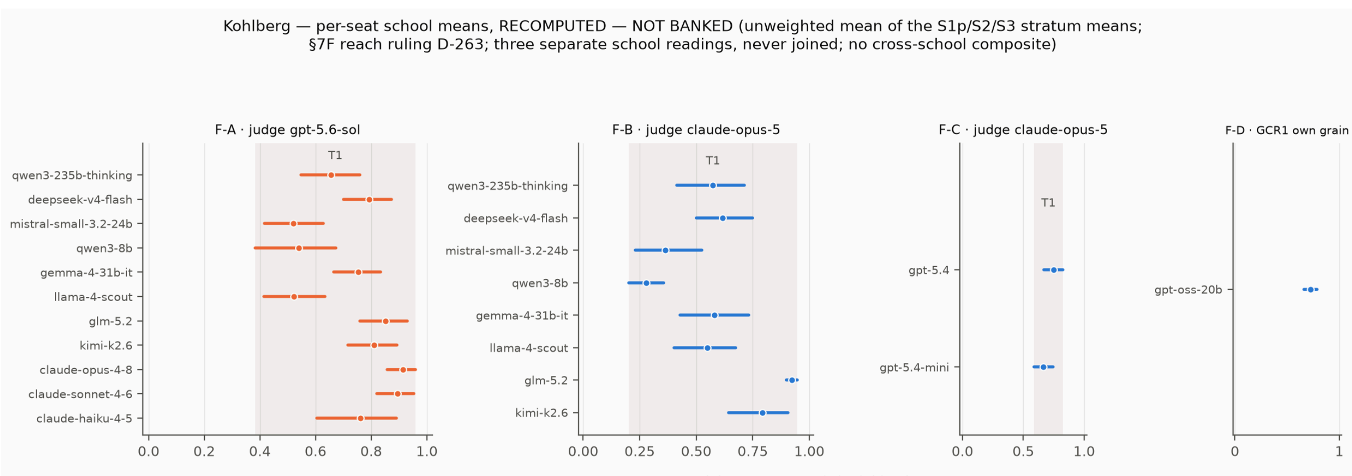

4.3 CI-backed tiers per same-judge frame

frame	setting	tier	members (census order; no within-tier order)
F-A	S2	T1	`qwen3-235b-thinking`, `deepseek-v4-flash`, `mistral-small-3.2-24b`, `qwen3-8b`, `gemma-4-31b-it`, `llama-4-scout`, `glm-5.2`, `kimi-k2.6`, `claude-opus-4-8`, `claude-sonnet-4-6`, `claude-haiku-4-5`
F-A	S3	T1	`qwen3-235b-thinking`, `deepseek-v4-flash`, `mistral-small-3.2-24b`, `qwen3-8b`, `gemma-4-31b-it`, `llama-4-scout`, `glm-5.2`, `kimi-k2.6`, `claude-opus-4-8`, `claude-sonnet-4-6`, `claude-haiku-4-5`
F-A	S1	T1	`qwen3-235b-thinking`, `deepseek-v4-flash`, `mistral-small-3.2-24b`, `qwen3-8b`, `gemma-4-31b-it`, `llama-4-scout`, `glm-5.2`, `kimi-`

			`k2.6`, `claude-opus-4-8`, `claude-sonnet-4-6`, `claude-haiku-4-5`
F-B	S2	T1	`qwen3-235b-thinking`, `deepseek-v4-flash`, `mistral-small-3.2-24b`, `qwen3-8b`, `gemma-4-31b-it`, `llama-4-scout`, `glm-5.2`, `kimi-k2.6`
F-B	S3	T1	`qwen3-235b-thinking`, `deepseek-v4-flash`, `mistral-small-3.2-24b`, `qwen3-8b`, `gemma-4-31b-it`, `llama-4-scout`, `glm-5.2`, `kimi-k2.6`
F-B	S1	T1	`qwen3-235b-thinking`, `deepseek-v4-flash`, `mistral-small-3.2-24b`, `qwen3-8b`, `gemma-4-31b-it`, `llama-4-scout`, `glm-5.2`, `kimi-k2.6`
F-C	S2	T1	`gpt-5.4`, `gpt-5.4-mini`
F-C	S3	T1	`gpt-5.4`, `gpt-5.4-mini`
F-C	S1	T1	`gpt-5.4`, `gpt-5.4-mini`

As in §2.3 and §3.3, **every Kohlberg panel yields the single tier T1** under the frozen rule — a chained band, not an equivalence statement. Kohlberg's chains differ in where they are anchored: unlike Schwartz, **this school's S1p column is not saturated** — it is the board's widest SMALL stratum, reaching from a FLOOR cell (`qwen3-8b`, 0.000, under the `claude-opus-5` leg) to the high-0.8 range within these frames — so even the S1p panels chain across nearly the whole scale (with the SMALL-stratum caveat at full strength: n = 4, one cell moves a mean visibly). The anti-conservativeness rider (NIR R1/R8, including a declared template class on the koh S3 slice) applies as before.

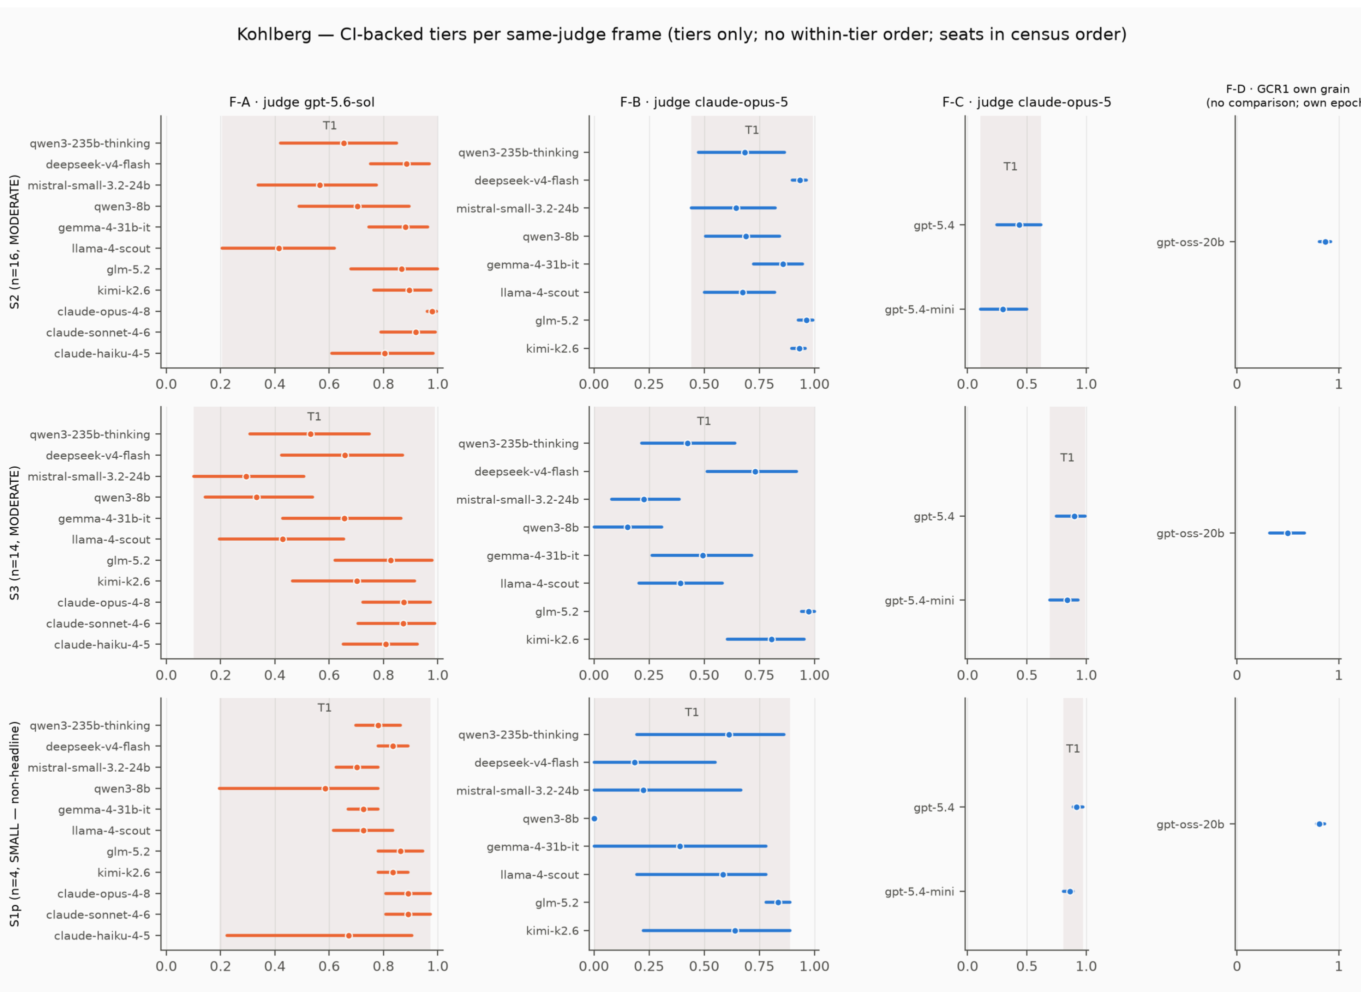

4.4 What the Kohlberg board shows (description at tier grain)

S1p (n = 4, SMALL) spreads instead of saturating. F-B spans 0.000–0.835; F-A spans 0.585–0.890 with no exact-1.000 cell anywhere in the school (Kohlberg contributes none of the board's eleven ceiling cells). The school's fresh-item floor stratum is doing discrimination-shaped work that the other two schools' saturated S1p columns do not — but it stays SMALL and non-headline, and its apparent spread is exactly where the one-cell-moves-the-mean caveat binds hardest.

S2 (n = 16) sits high in most rows. F-A spans 0.413–0.979; F-B spans 0.643–0.963. The re-administered B3 kohlberg S2 cells (FC-057) sit inside these bands.

S3 (n = 14) carries the school's spread in both main frames — F-A 0.294–0.874, F-B 0.151–0.972 — again quasi-continuous, which is the chaining geometry. In F-C the two OpenAI-substrate seats show that substrate's banked inverted setting order at its most visible: their S3 cells (0.900, 0.838) sit far above their S2 cells (0.434, 0.295) — a shape statement about that leg, never a ranking.

The GCR1 seat, read beside the frames: `gpt-oss-20b`'s Kohlberg cells are 0.808 (S1p) / 0.866 (S2) / 0.498 (S3), Kohlberg school mean 0.724 [0.659, 0.785] (`RECOMPUTED — NOT BANKED`, recipe §1.6), own grain, own epoch, single-OOF. No comparative statement attaches.

Tier-grain summary. Within any one frame, no Kohlberg CI-separation between any seats at any setting: one chained band per panel, with an unsaturated S1p column and the S3 column carrying the widest spread.

---

§5 Bars and exclusions

5.1 The Family-P seat-comparison bar (cited, never approximated)

Family-P (the profile family) contributes **nothing** to this report: firewall **G2** is absolute — no Family-P scored quantity appears in any table or sentence here. Beyond G2, the whole-family standing bars named by the commissioning boot bar a Family-P seat comparison outright, by citation (INV §7·C/§7·D; ids per the banked bars register):

- **BAR-INV-7C-1** — *"Family-P whole-family standing bars (absolute, not merely pre-Gate-M): no ranking/ordering/tier/composite/best-worst, including in prose."*
- **BAR-INV-7C-4** — *"No cross-stratum or cross-substrate comparison, pooling, averaging or differencing in either direction (the A-10 fence, absolute)."*
- **BAR-INV-7C-13** — *"Per-seat coder-agreement cells are not printed — agreement is a property of the MEASUREMENT INSTRUMENT, not of any model …"* (quoted in part)
- **BAR-INV-7D** — the H5 roster-grain dispersion statistic is **NOT COMPUTED**; the verdict stays NOT-EVALUABLE and is never converted into a number.

Family-P seat surfaces are absolute (per-seat descriptive only, no comparison across seats); that is why this commission compares seats on the three schools' **Family-C** boards only — all three schools have Family-C boards, so every school reading here is a competence reading.

5.2 The two absent seats (declared, with riders)

- **`google/gemma-3n-e4b-it` — PROVIDER-ABSENT-AT-REPROBE (2026-08-28).** The fresh fail-closed preflight ($0.00, 0 paid cells) found the slug does not resolve and no gemma-3n variant remains on the provider; any substitution would be an identity-swap roster amendment, never a dated re-pin. Consequence of record: the 15/16 path is unreachable as designed; the board ceiling on the current roster is **14/16 — attained** (ADD §4).

- **`openai/gpt-5.5` — EXCLUDED per amendment A-22** (REDUCED_ROSTER / TRANSPORT_FAIL). The seat renders at **denominator 0** and carries no result; it is not dropped from the sealed roster — registering is not dropping. Status narration carries the dated three-step transport sequence of record: the two bounded 2026-08-07/08 attempts with 3/3 smoke episodes prose-only (tool-non-engaging); the MCP bisection **measured 2026-08-08** (declaration banked 2026-08-19) — discovery ≠ usable transport, 9/9 mock calls failed; and the fresh 2026-08-28 one-call re-check, 0 ok / 1 fail, verdict TRANSPORT_FAIL (ADD §4). No estimate, proxy, or reconstruction is offered for either absent seat.

5.3 The cross-school composite bar

The §7F bar on a per-seat mean across schools is **absolute** for this report: the three school readings are never joined into one number, one axis, or one ordering. *"There is no composite 'moral psychology score'"* (CS §4); the D-263 reach ruling licenses **within**-school means only. The computation module encodes the bar: its sole cross-school function raises.

5.4 The within-tier-order bar (pre-Gate-M)

Any within-tier ordering is BARRED until Gate M. Tier members print in census order everywhere, including inside figures; tier labels order **tiers**, never members.

5.5 Capability-scalar absence

No capability covariate is available on any surface of this report (the BAR-STAND-1 / RG-STATUS-7 pattern): the roster's published capability scalars are pinned later via the sealed snapshot rule, and **no capability-adjusted reading of any tier is possible here**. Tier structure must not be read as a capability ordering, and no such adjustment is computed.

5.6 What else does not appear here

No confirmatory H-claim (the preregistered hypotheses remain unexecuted on this substrate as of the 2026-08-19 dated statement; QUARANTINED-DESCRIPTIVE thereafter); no judge-validity claim (the judge annex is consistency/diagnostic only); no post-A-14 measured or re-judged result (none exists); no pooled cross-epoch quantity; no per-seat results for the 160 never-administered expansion items; no banked B1/B2/frontier grand mean (the §B.6 item-3 bar; the recomputed aggregates of INV §C.2.6(c) served as check targets only and are not reprinted); no per-seat cross-school "overall" mean (§7F, ADJ-3: where a bar's reach is unruled, ambiguity resolves toward exclusion).

---

§6 Provenance, verification, and relations

Sources of record. The §1.1 substrate table, pins re-verified at session open (all 8 MATCH in both working copies). Board cells are copied from INV §C.2.1–.4 (FC-013…FC-033) and ADD §3; every new quantity is recomputed from the RAW judge legs named in §1.3 via `engine/school_comparison.py` (tests: `engine/tests/test_school_comparison.py`, 24/24 green at build). Verification battery: `reports/figures/20260829-school-comparison/CHECK_LOG.txt` (positive control 32/32; banked cells 198/198; secondary aggregates as check targets; tier-rule gate one line per panel; STOP-gate exits non-zero on any FAIL).

Declared no-verdicts on this surface. `sch-s2-003` × `gpt-5.4` (schwartz S2, judge safeguard refusal — stratum n = 15, the lane's only board no-verdict, FC-055) and `sch-s3-013` × `gpt-oss-20b` (schwartz S3, subject-side `empty_output`, DECLARED — stratum n = 13 of 14, ADD §3). Disclosed, never imputed.

Relations. This report **extends** the FC results report (FCR, non-computing citation surface) with the first computed tier/CI reading of the 14-seat state; it edits no earlier edition and supersedes nothing. A sibling commission banked the same day (MPD-309/D-264, the model-comparison paper boot) overlaps on the Family-C school grain; per its banking reconciliation, **neither commission rides, widens, or supersedes the other**, and the overlap is disclosed here by pointer only.

Open tensions carried, not resolved. CONFLICT-A (`llama-4-scout` schwartz S3 leg 2: 0.563 edition-of-record vs .562 in three readout renders — exact mean 0.5625; the banked TENSION-1 rounding divergence) and CONFLICT-B (`gpt-5.4` kohlberg S1p: 0.918 of record vs .917 — exact mean 0.9175, ADJ-2) ride the affected cells below; this report prints the edition-of-record values only. Two further divergences of the same exact-half render class surfaced in this session's recompute and are disclosed, not repaired: `claude-haiku-4-5`'s F-A mft S3 (banked 0.662; exact mean 0.6625; this recompute's ROUND_HALF_UP render 0.663) and its F-A kohlberg S1p (banked 0.672; exact mean 0.6725; recompute render 0.673) — both inside the banked ±0.0005 envelope, and the banked values print throughout this report. The B2 `claude-opus-5` per-seat surface carries its banked provenance tension (FC-036: no published readout carries it; primary source is the leg record tree via the verified derived-census recomputation) — carried, never resolved.

Rendered 2026-08-29 by the cross-LLM school-comparison session (MPD-308/D-263 commission). $0 — no paid API call of any size. Gate M remains the sole validity path; nothing here is validity evidence.