

Administration Is Part of the Construct

Measuring Moral-Psychology Schools in Language Models — preregistered report, v2


Anonymous Author(s)

Gate M: PENDING · every result-shaped quantity is quarantined-provisional · no expert ratings exist as of this writing

The contribution: a conjunction, three constructs, one firewall

No existing benchmark operationalizes all three moral-psychology schools — Moral Foundations Theory, Schwartz Basic Human Values, and Kohlberg and neo-Kohlbergian moral development — as separate theory-native constructs, and none applies the competence-versus-default-profile firewall to them. The contribution is that conjunction, under measurement-disciplined conditions.

We claim no first on any ingredient (§1.2).



Three schools, three constructs, two families across Firewall G2.

- MFT** → **intuition-conflict weighting**: which concern is decisive when foundations conflict.
- Schwartz** → **circumplex-respecting tradeoffs**: recovery per administration condition.
- Neo-Kohlberg** → **justification-structure scoring**: never a stage label for a model.

Sealed instrument: **406 public items paired one-to-one with 406 withheld shadow twins** — **812 item surfaces** — **352 sealed rubrics**.

Default profiles are **never ranked, never tiered, and never reduced to a single number** — a permanent prohibition, reopen condition "never" (PM1-C095), held by five enforcement points and a test lock.

Administration is part of the construct: one sample per (item, variant, administration), **never averaged** before it is reported; regime separation fail-closed before any model call; seed in the sample identity.

B2-01

Competence family — per-seat, per-leg surfaces

Thirteen seats, four disjoint substrates, never merged. Per seat **120 subject cells, 102 judged** (4 / 16 / 14 per school); dual-judged **204 verdicts**, single-judged **102**; denominators **360 / 360 / 240 / 600**; **1 disclosed no-verdict** (frontier leg, schwartz-S2, **n = 15**, never imputed). "Frontier" names a vendor's flagship line, never a measured tier.

Table 1 — per-seat, per-leg school × setting surfaces; legs never averaged, nothing merged across substrates. Rows in census roster order, which is not a score order, and no row from one substrate is comparable with a row from another.

substrate · leg · seat	koh S1p	koh S2	koh S3	mft S1p	mft S2	mft S3	sch S1p	sch S2	sch S3
B1 · leg 1 · claude-opus-4-8	0.890	0.979	0.874	1.000	0.966	0.682	1.000	0.817	0.386
B1 · leg 1 · claude-sonnet-4-6	0.890	0.920	0.873	0.980	0.972	0.822	0.500	0.938	0.891
B1 · leg 1 · claude-haiku-4-5	0.672	0.804	0.808	0.980	0.960	0.662	0.750	0.909	0.643
B2 · leg 1 · gemma-4-31b-it	0.390	0.857	0.494	0.955	0.846	0.177	1.000	0.861	0.516
B2 · leg 1 · llama-4-scout	0.585	0.673	0.391	0.485	0.386	0.206	0.919	0.867	0.464
B2 · leg 1 · glm-5.2	0.835	0.963	0.972	0.955	0.988	0.684	0.750	0.964	0.686
B2 · leg 2 · gemma-4-31b-it	0.725	0.881	0.656	0.975	0.900	0.088	1.000	0.880	0.500
B2 · leg 2 · llama-4-scout	0.725	0.413	0.429	0.930	0.813	0.366	1.000	0.902	0.563
B2 · leg 2 · glm-5.2	0.863	0.867	0.826	0.980	0.978	0.621	1.000	0.923	0.871
frontier · leg 1 (v2) · gpt-5.4	0.918	0.434	0.900	0.955	0.463	0.516	1.000	0.565	0.582
frontier · leg 1 (v2) · gpt-5.4-mini	0.863	0.295	0.838	0.955	0.355	0.401	1.000	0.466	0.505
B3 · leg 1 · qwen3-8b	0.000	0.688	0.151	0.480	0.738	0.000	0.500	0.673	0.257
B3 · leg 1 · mistral-small-3.2-24b	0.223	0.643	0.225	0.895	0.778	0.000	0.725	0.759	0.346
B3 · leg 1 · kimi-k2.6	0.640	0.930	0.806	0.730	0.889	0.476	0.750	0.944	0.473
B3 · leg 1 · deepseek-v4-flash	0.184	0.933	0.731	0.938	0.799	0.226	1.000	0.947	0.329
B3 · leg 1 · qwen3-235b-thinking	0.613	0.684	0.425	0.725	0.625	0.133	0.719	0.709	0.089
B3 · leg 2 · qwen3-8b	0.585	0.703	0.331	0.730	0.914	0.000	0.750	0.842	0.321
B3 · leg 2 · mistral-small-3.2-24b	0.703	0.566	0.294	0.945	0.791	0.039	0.750	0.833	0.468
B3 · leg 2 · kimi-k2.6	0.835	0.895	0.701	0.750	0.981	0.477	1.000	0.936	0.379
B3 · leg 2 · deepseek-v4-flash	0.835	0.885	0.657	0.963	0.902	0.318	1.000	0.933	0.491
B3 · leg 2 · qwen3-235b-thinking	0.780	0.653	0.531	0.750	0.845	0.131	0.500	0.789	0.093

Bounded interpretation, fixed in advance: this evidence **supports the measurement architecture, not stable ordering** (PM1-C050). It establishes nothing about which model is better on any construct.

B2-02, B2-03, B2-04, B2-08, B2-25

Profile family — the sealed bundle-floor gate

The value-priority construct produces **no per-seat priority profile anywhere in this program**, and that absence is a design fact rather than attrition. The construct is two-stage and strictly ordered: a structure-recovery gate must be evaluated before any priority vector may be reported.

The unit of that gate is a pseudo-respondent bundle. The sealed minimum is **30 bundles per seat** — a constant anchored at 30 *respondent* bundles in the pre-freeze power simulation. The administration frame yields **3 bundles per seat**, and the preregistered **DECLINE fired for every seat: 18** recorded seat-gates (**5 + 5 + 8**) on the three t = 0 substrates of the comprehensive report, and by construction on the later extension arc's three seats.


Table 2 — the sealed bundle-floor gate. A gate outcome per seat, never a score and never an ordering; no seat is named and no construct-vector quantity appears.

sealed gate unit	sealed minimum per seat	administration frame yields	seat-gates evaluated (comprehensive 13-seat report)	DECLINE
pseudo-respondent bundle	30 respondent bundles	3 bundles	18 (5 + 5 + 8)	18

One later arm reached the floor, at exactly **30 bundles per seat**, and the gate **PERMITTED** on all five of its seats — whereupon the sealed recovery returned **INDETERMINATE** on every one, because **two of its three criteria have no sealed implementation**: the three-way verdict ladder binding on real data, not a bug, not a null result, and not a pass. Non-recovery is never read as a model disposition, and that arm's priority vectors and correlation matrices are never rendered. A sealed human-derived constant — a floor stated in respondent bundles — does not survive the port to a language model, because the frame that generates bundles has no respondents in it.

B2-19, B2-20

Governance, the gate, and limitations



Fail-closed claims-ledger build and the census-clean identity.

- Fail-closed ledger build** — the build **fails rather than emitting an unbound number**.
- DECLINE** is a preregistered outcome; INDETERMINATE-via-NOT_EVALUABLE is a pre-declared indeterminacy, not a pass. No public-versus-withheld performance comparison appears anywhere (PM1-C096).

Gate M: PENDING — the judge must be validated **as an instrument**, against a sealed floor, before it grades a headline. No expert ratings exist, and no expert-validated judge-performance claim is made anywhere (PM1-C093).

Instrument vintage. Every score in this section was produced under the instrument's prior 2026-08-08 seal and the pre-revision judge prompt; the 2026-08-09 re-seal of record post-dates every reported measurement — it changed no counts, but it revised judge-prompt rendering and nineteen item surfaces, so no number here was produced under, or is comparable to output of, the current sealed instrument.

- No validated judge**: current evidence is **feasibility evidence**, and agreement between two unvalidated instruments is never validity.
- Two separate boards, stated separately and never compared with each other** (Firewall G2); **not a representative sample** of any model population; "frontier" names a vendor's flagship line, never a measured competence tier.
- Ranking, tiers, hypothesis verdicts and generalization are all limited — **prohibited rather than merely hedged** (PM1-C090...C094).

Evidence banner — mandatory; travels with every quantity on this poster.

- All results in §8 are **descriptive, exploratory, pre-Gate-M** results on the program's measured boards, and they inherit the render law of the reports they bind to.
- Every score is the output of an **unvalidated** judge or coder, or an agreement between two unvalidated ones; **validity is Gate-M-only**, Gate M has not run, and agreement between two unvalidated instruments is never validity.
- Where a substrate is dual-judged there is **one table per judge**, and the two legs are **never averaged into a single number**; where a substrate is single-judged, its surface is one unvalidated judge's output alone.
- Substrates are **never merged**: coverage grew by adding substrates, not by growing one, and verdict denominators differ across them. Seat rows appear in their declared roster order, which is **never a score order**.
- Nothing here establishes a ranking (PM1-C090), a competence tier (PM1-C091), a hypothesis verdict (PM1-C092), judge validity (PM1-C093), or generalization beyond these seats (PM1-C094).

Colophon — v2. Evidence pin: commit 3261ddd9 . Seal of record: sealed instrument manifest, digest be0a11a3a4d7 . Claim-audit rows displayed on this poster — instrument: B2-01 · competence family: B2-02, B2-03, B2-04, B2-08, B2-25 · profile family: B2-19, B2-20. Gate M: **PENDING**; quarantined-provisional throughout. *Template / venue TBC*.