

PROVISIONAL · DESCRIPTIVE · EXPLORATORY — judge validity is NOT established and Gate M has not run

Every score on this sheet is the output of an **UNVALIDATED** judge — `claude-opus-5` and/or `gpt-5.6-sol`, both pre-Gate-M — or an agreement statistic between two unvalidated judges. **Consistency is not validity. Judge validity is Gate-M-only, and Gate M has not run.** Therefore: **no rankings, no ordering of seats by score, no best/worst language, no cross-substrate merge.** Seats appear in the census's declared roster order everywhere, never in score order. These are **within-theory COMPETENCE** results (faithful reasoning on request); competence-under-instruction is not general capability and nothing here is a capability claim. **Firewall G2** holds throughout: no competence number shares a table, analysis, or sentence with a default-value profile number.

src: FC-M render law L8-12 · FC-R A-1 render law L11-46 · Scope L50-67

What Family-C measures — and the two channels underneath it

Family C asks one question of each seat: **when explicitly asked to reason inside a named moral-psychological school, does the model reason faithfully within it?** That is not the Family-P question (which school's verdicts *unprompted* behaviour resembles), and the two never share an analysis, table, or sentence.

Three schools, three different measurement objects

- kohlberg — Kohlberg's cognitive-developmental theory in its **neo-Kohlbergian** form: reasoning-structure sophistication.
- mft — Moral Foundations Theory: foundation-faithful reasoning.
- schwartz — Schwartz's **refined 19-value** model: value-location and value-conflict reasoning.

Four settings, two channels — the split is load-bearing

Every seat receives the same **120 subject cells** = S1a 18 · S1p 12 · S2 48 · S3 42. The **judged** geometry is **102 cells per seat** = per school S1p 4 · S2 16 · S3 14.

- Mechanical channel.** The **18 S1a cells per seat** are fixed-key, mechanically scored, and **never sent to a judge.** They are carried inside every substrate census (inside the 360/360, 240/240 and 600/600 counts) and in the cost model, but sit **OUTSIDE every judged surface** — **no mechanical-channel quantity is published for any of the 13 seats**, and none should be inferred.

- Judged channel.** S1p is free production graded against a rubric; **S2** is an agentic single episode inside a fictional organisation's charter in a mock MCP tool world under incentive or framing pressure, with the rubric scoring *honest behaviour*; **S3** is multi-turn, the position probed by new facts and by bare pressure to change the answer.

The single most load-bearing design fact

S2 is scored on the POST-EPISEDE DEBRIEF — the reasons channel — not on the episode's final message. S2 episodes end via a terminal tool call, so `final_text` is empty **BY DESIGN**: 144/144 census-clean judged B2 S2 records carry an empty `final_text`. Two consequences cost the programme real work: a route that fails to capture the debrief yields an **empty-but-scored** channel — a genuine, non-null score of about 0.0 that a `score != null` census cannot distinguish from a real low score; and a smoke gate demanding non-empty `final_text` on S2 is an **instrument error** that would fail seats for behaving as designed.

src: FC-R §0 L69-76 · §4.1 L507-516 · FC-M §0.1-0.3 L55-65 · §2.1 L161-172 · §2.3 L178-188 · COMP §2.1 L290-314 · §2.2 L316-346

The measured board — 13 seats, 4 substrates, per-seat denominators

Family-C measured coverage: 13 seats · 4 distinct substrates (never merged) · per-seat denominators. Dual-OOF seats carry 204 verdicts (102 judged cells × 2 judges); single-OOF seats carry 102 (LOO firewall). frontier gpt-5.4 carries 1 disclosed no-verdict (schwartz-S2 over n=15). The declared exclusion is rendered, not omitted: gpt-5.5 is a REDUCED-ROSTER seat at denominator 0, never administered, and is not among the 13. All judges UNVALIDATED pre-Gate-M; this is a coverage map, not a board — no scores appear on it.

substrate	seat	subject cells	judged cells	judge legs	verdicts	OOF class
B1	claude-opus-4-8	120	102	gpt-5.6-sol	102	single-OOF
B1	claude-sonnet-4-6	120	102	gpt-5.6-sol	102	single-OOF
B1	claude-haiku-4-5	120	102	gpt-5.6-sol	102	single-OOF
B2	gemma-4-31b-it	120	102	claude-opus-5 + gpt-5.6-sol	102×2 = 204	DUAL-OOF
B2	llama-4-scout	120	102	claude-opus-5 + gpt-5.6-sol	102×2 = 204	DUAL-OOF
B2	glm-5.2	120	102	claude-opus-5 + gpt-5.6-sol	102×2 = 204	DUAL-OOF
frontier	gpt-5.4	120	102	claude-opus-5 (v2)	102 (1 NV)	single-OOF
frontier	gpt-5.4-mini	120	102	claude-opus-5 (v2)	102	single-OOF
B3	qwen3-8b	120	102	claude-opus-5 + gpt-5.6-sol	102×2 = 204	DUAL-OOF
B3	mistral-small-3.2-24b	120	102	claude-opus-5 + gpt-5.6-sol	102×2 = 204	DUAL-OOF
B3	kimi-k2.6	120	102	claude-opus-5 + gpt-5.6-sol	102×2 = 204	DUAL-OOF
B3	deepseek-v4-flash	120	102	claude-opus-5 + gpt-5.6-sol	102×2 = 204	DUAL-OOF
B3	qwen3-235b-thinking	120	102	claude-opus-5 + gpt-5.6-sol	102×2 = 204	DUAL-OOF
EXCLUDED	gpt-5.5	0	0	— (never administered)	0	REDUCED-ROSTER

Figure 9 — Family-C measured coverage: 13 seats across 4 distinct substrates (never merged), with per-seat denominators. Rows are in census roster order within each substrate band; the row order is not an ordering of seats by anything. Every seat received 120 subject cells of which 102 are judged. Dual-OOF seats (B2, B3) carry 204 verdicts (102 judged cells × 2 independent judge legs); single-OOF seats (B1, frontier) carry 102 under the sealed leave-one-out firewall. frontier / gpt-5.4 carries 1 disclosed no-verdict (schwartz-S2 over n=15, never imputed). Judges named per band: B1 gpt-5.6-sol; B2 claude-opus-5 + gpt-5.6-sol; frontier claude-opus-5 (v2) corrected leg; B3 claude-opus-5 + gpt-5.6-sol. ALL JUDGES UNVALIDATED pre-Gate-M.

src: FC-R §0 L98-107 (figure + caption) · COMP §2.3 L397-405

Per-seat score shapes — one panel per seat, grouped by substrate




Figure 5 — 13 measured seats, per-seat score shape across the nine school × setting strata, one panel per seat, grouped by substrate. Panels sit in census roster order within each substrate band, and the bands are B1 → B2 → frontier → B3 in substrate order, NOT in any order of score. The substrate rows differ in judge stack: B1 is judge gpt-5.6-sol only (SINGLE-OOF, 102 verdicts/seat); B2 shows both legs (claude-opus-5 and gpt-5.6-sol, DUAL-OOF, 204 verdicts/seat); frontier is judge claude-opus-5 only (SINGLE-OOF, 204 verdicts/seat). The four bands are four distinct substrates and are never merged — a panel in one band is not comparable to a panel in another as a measurement, only as a shape. All judges UNVALIDATED pre-Gate-M. All S1 points are SMALL strata (n=4).

src: FC-R §2 L362-373 (figure + caption)

Score shapes — labels first, no orderings

- Every stratum is labelled before any observation is drawn from it.** SMALL: every S1 cell is n=4 per school per seat (consistency S1 strata are n=12 on B2, n=20 on B3); MODERATE: S2 n=16, S3 n=14 per school per seat. AGGREGATE: pooled within one leg, plus FLOOR and CEILING.
- The modal shape is S2 above S3 within a seat, same school: 8 of 9 school × seat comparisons on B1 (one judge), 8 of 9 on each B2 leg, 15 of 15 on each B3 leg.**
- Exactly three exceptions** on those three substrates, all in kohlberg and all near-ties: B1 `claude-haiku-4-5` (.804 versus .808, judge `gpt-5.6-sol`); B2 `glm-5.2` on leg 1 (.963 versus .972); B2 `llama-4-scout` on leg 2 (.413 versus .429). Reported rather than rounded away.
- The frontier band INVERTS that shape:** both frontier panels put S2 below both S1 and S3 in all three schools (judge `claude-opus-5`, v2 corrected, single-OOF). This is a within-seat contrast between two settings — never an ordering of seats.
- CEILING:** cells at exactly **1.000** are almost exclusively `sch-S1`, plus one mft-S1 cell at B1 `claude-opus-4-8` — every one of them n=4, i.e. **saturated-on-small-n**. FLOOR: exact .000 mft-S3 cells on B3 (both legs); B2 `gemma-4-31b-it` mft-S3 .088 (leg 2 `gpt-5.6-sol`, n=14; leg 1 reads .177).
- Floor and ceiling strata are EQUALLY uninformative for discrimination.** Both are disclosed as **design inputs, never as seat verdicts** — a stratum that is a point mass tells you about the instrument, not about the seats sitting in it.
- Cell values above are the render-time recomputation from the per-cell judge records, verified against the published readouts **where one exists** — the B2 `claude-opus-5` per-seat surface has no published readout; the leg record tree is its only primary source (FC-R §9.2c).

src: FC-R §0.1 L109-123 · §1 L131-134 · §2 L337-377 · §1.3 L245-248 · §8 items 4-5 L679-686

The judge pair — consistency is NOT validity

- Only **B2** and **B3** carry a consistency surface. **B1** and the frontier substrate are single-OOF and carry **NONE** — there is no second reading of those cells and no way to bound judge variance on them. **B2** and **B3** are **NEVER pooled**. Delta convention on both: `gpt-5.6-sol` minus `claude-opus-5` on the same cell.

statistic	B2 · 306 paired cells, 3 seats	B3 · 510 paired cells, 5 seats
mean abs Δ	0.1618	0.1476
exact agreement	0.471	0.537
within-0.1 (epsilon)	0.670	0.749
signed mean Δ	+0.0403	+0.0599
no-verdicts (final)	0	0

- One convention, named.** Within-0.1 is shown under the **epsilon rule** for both substrates. The float-iter counts are .667 (B2) and .741 (B3); the two published readouts headline *different* conventions, and quoting them side by side overstates the difference by the boundary cells alone.
- Same direction on both substrates** — `gpt-5.6-sol` reads higher on average — and larger on B3. The B3 offset **+0.0599** equals the difference of the two legs' published grand means (.628 — .568) to four decimals: an **arithmetic cross-check, NOT a merge**, not a school effect, and not a correctness statement about either judge.
- The highest agreement on the board is partly a FLOOR ARTIFACT** (FC-R §5.2): B3 mft-S3 (mean abs Δ .030) is in part two judges agreeing that a cell sits at the bottom of the scale. Cheap agreement is not judge quality.
- Every consistency figure is a property of the judge pair **under the PRE-POLES prompt regime**, never a property of a seat.
- Table abridged from FC-R Table 5.3: the *within-0.25 row* — the one row where the B2/B3 direction reverses — is not reproduced, and §5.3's *guard travels with what is: the B3-favouring rows are "not a claim the judges got better"*.

src: FC-R §5 L596-609 · §5.0 L611-625 · §5.1 L629-657 · §5.2 L675-708 · §5.3 L723-742 · §3.1 L429-434

Dual-judge consistency on B3 — a property of the judge pair, never of a seat

B3 (510 pairs) dual-judge consistency — a judge-pair property, never a seat property consistency ≠ validity (both judges UNVALIDATED pre-Gate-M)




Figure 7 — B3 dual-judge consistency: the absolute-delta distribution (left) and mean absolute delta by school × setting stratum (right). Judges claude-opus-5 and gpt-5.6-sol, DUAL-OOF, 510 paired cells (5 seats × 102); the delta convention is `gpt-5.6-sol` minus `claude-opus-5` on the same cell. Per-stratum n is labelled on each bar; the S1 strata are SMALL (n=20). This is a judge-pair property, never a seat property. CONSISTENCY IS NOT VALIDITY — both judges UNVALIDATED pre-Gate-M. Within-0.1 is shown under the EPSILON rule (.749), the float-iter count is .741). Every figure here is a property of the pair under the PRE-POLES prompt regime.

src: FC-R §5.2 L713-719 (figure + caption) · §5.0 L611-625

LIMITATIONS — these bind every number on this sheet

1 · The judges are UNVALIDATED

Every number here is one unvalidated judge's output, or an agreement statistic between two unvalidated judges. **Gate M — expert consensus labels — is the sole validity path and has not run.** No amount of additional coverage changes that: **B3 added 5 seats and 1,020 verdicts and moved validity by exactly zero.** Consistency is not validity; single-OOF substrates carry no consistency surface at all; the four substrates are never merged.

12 · Every score is a PRE-POLES grade

On **2026-08-09**, after every leg here had been judged, amendment **A-12 / W1** made the judge prompt render each criterion's **V8 pole anchors** and re-versioned the rubric fingerprint (`rubric-fp=v2-poles`). **Pre- and post-2026-08-09 judge results are NOT comparable**, and every pre-fix judge cache is retired by design, so any repeat is a new measurement rather than a reproduction. **No number in the source report changes.**

13 · No individual grade can be spot-checked (as annotated)

The shipped judge prompt does not implement the **binding judge-prompt specification: no evidence span was ever requested or stored anywhere on the reasons channel, so no individual grade on these surfaces can be spot-checked against the text that produced it.** Read "sealed" as **BINDING** — **no hash of the judge template exists in prereg/SEALED**; / the freeze sealed pools, rubrics, shadows and the prereg constants only, so there was never a sealed template for the shipped builder to be diffed against. **No number moves.**

src: FC-R §8 items 1-2 L868-875 · item 12 L919-934 · item 13 L962-1006 + MPD-189 annotation L1008-1022

The retracted frontier artifact — and the road forward

A worked example: a headline that was wrong, and was withdrawn

- The frontier arc's original headline — both OpenAI seats scoring near zero on the agentic setting — was **WRONG and was withdrawn**.
- Root cause:** the subject route failed to capture the debrief. `debrief.answers` == [None, None, None] for **48/48 S2 cells on both seats**, against **0/48 empty** on every OpenRouter and Claude-CLI route over the identical items. A blind judge then faithfully scored an **empty channel** at about 0.0.
- A score != null, census CANNOT detect an empty-but-scored channel** — an empty debrief still yields a real, non-null score. That is the disclosed integrity gap.
- Fixed, re-run, re-judged:** 96/96 census-clean with 48/48 populated debriefs per seat, on a v2 leg that supersedes v1 for the whole frontier seat set. Engagement was never the issue — tool use was **16/16** (`gpt-5.4`) and **14/16** (`gpt-5.4-mini`) on the spot-checked stratum.
- The v1 column exists ONLY as a labelled retracted artifact** and must never be quoted as a result. The corrective — a reasons-channel completeness check beside `score != null` — ran at full-run scale on B3: gate **PASS 240/240** debriefs populated, 0 empty, 5/5 seats.

The road forward, in the lane's own order

(i) wave-D authoring completes the ratified quota (**15 of 16 pairs remaining**, merge-gated on fidelity review and the V10 scorer round-trip) [2026-08-11 review annotation: overtaken same-day — all 15 reviewed, Wave D MERGED + RE-SEALED, MPD-197/D-174, E-4; historical as of the 2026-08-10 assembly date]; (ii) the judge prompt-builder repair executes **BEFORE** Gate M; (iii) the **Gate-M vehicle question** is ruled and the expert arm runs; (iv) **only then** is any confirmatory sweep or headline claim possible. Baseline debts against this path are **disclosed and UNAUTHORIZED**: wave-D 1-0 **≈\$0.55-\$0.60** · W5 76-cell re-baseline **\$0.70** · re-score of the exploratory caches **\$0** cash on incumbent OAUTH routes, metered fallback **\$7-\$82** by judge class.

src: FC-R §1.3.1 L257-299 · §4.4 L567-582 · FC-M §7.1 L354-366 · COMP §5 L705-1713

Gate M is the sole path from these descriptive surfaces to any confirmatory claim — and it has not run

Internal owner briefing, assembled from banked Family-C records. **\$0 session — no model or API calls were made in producing this sheet, and no new statistic was computed anywhere on it;** every figure is a banked local render. **No benchmark item text appears — item identifiers only.** Firewall G2: no Family-P quantity appears on this sheet. Nothing here is a ranking, an ordering, a best/worst statement, or a capability claim.

src: A-1 render law L11-46 · the \$0/no-new-statistic sentence is a build-session fact (MPD-198/D-175), not an FC-R claim