

NEVER RANKED — absolute, and by construction, not merely pre-Gate-M


MP-v2 Family-P (profiling) — internal owner briefing, 2026-08-10. Every coded number on this sheet is the output of an **UNVALIDATED** coder — leg A `claude-opus-5` or leg B `gpt-5.6-sol` — or a comparison between two of them. **Coder agreement is not validity**: two coders agreeing is a property of the coder pair, and no quantity of seats, cells or legs converts one into the other. **Gate M is the sole path to coder validity, and Gate M has not run.**

None of the three constructs has a better/worse direction — an intuition-weight profile (MFT), a value-priority profile (Schwartz), a distribution of justification structures (Kohlberg). So there is **no ordering, tier, composite, best/worst or leading/trailing language about seats anywhere on this sheet** (A-1 scopes the fence to seats; channel-level evidence statements are permitted), and there never will be in Family P: passing Gate M would change what can be believed about the **CODERS**, and would not create a direction on the **CONSTRUCTS**. Row order is declared roster order and is never an ordering of seats by anything.

PROVISIONAL · DESCRIPTIVE · EXPLORATORY. Firewall G2: no Family-C competence quantity appears anywhere on this sheet. Content assembly only — **no new statistic, kappa is never pooled**. Seat names appear **only as figure panel labels, in declared roster order** — **never in a table or sentence**; and **no per-seat agreement, reliability or kappa quantity is attached to any seat anywhere in this program** (agreement is arc-level only, over denominators of no fewer than three seats).

MFT profiles — per seat, both coder legs, never averaged

Per-seat MFT foundation profile — share of the seat's own judged mft-P cells in which the coder marked the foundation **SALIENT** (raised / endorsed / argued_against). One panel per seat, grouped by substrate (row): **BOTH** coder legs are shown side by side and are never reduced to one code. Substrates are separate runs — never merged, never averaged, never ranked. Seats in declared roster order.



Per-seat MFT intuition-weight profiles, one panel per seat (18 panels), grouped into never-merged substrate rows. BOTH coder legs are drawn in every panel as two overlaid line series — leg A and leg B — never reduced to one code, never tie-broken, never averaged. Denominator: each seat's own 32 mft-P cells; foundations in sealed instrument order; seats in declared roster order over within their substrate row, which is not an ordering. Shares are per foundation on their own denominator and are never averaged into one MFT number. Both coders UNVALIDATED pre-Gate-M.

src: FP-R §1 Figure 2 · COMP §3.1 Figure 3a — *banked artifact, not re-rendered*

What profiling measures, and how it differs IN KIND

- Family P measures what **unprompted default behaviour expresses** — not whether a model can reason faithfully inside a named school when instructed.
- The difference is the **kind of object measured**. A competence surface has a **direction**: it can be scored and ordered. A **profile surface has no direction at all**.
- Moral psychology is descriptive**; no school defines a **"better" default profile**. A maintaining-norms justification is not a defect, a postconventional one is not an achievement, a value priority is not a score.
- The never-rank law is therefore **STRUCTURAL, not provisional** — a property of the constructs, not a caution awaiting validation.
- Dispersion is part of the P-MFT measurand and was not measured**: one variant per item at $t = 0$, so every surface here is **location-only**.

src: COMP §3.1 · FP-M §2.4 · FP-R A-1 render law

The instrument as administered — and the 74%


Per seat, identical on every $t = 0$ substrate: **38** kohberg-P + **32** mft-P cells (**dual-coded**) and **204** schwartz-P cells (**mechanical, never judged**) = **274 P cells, of which 70 are coded**.

- 204 of every seat's 274 P cells — 74% — never reach a coder at all.**
- Every agreement statistic therefore describes a **minority** of the administered profile surface. Reading *"the coders agreed at rate r "* as *"the Family-P data are reliable at rate r "* multiplies the claim by roughly four.
- The cell budget is **identical for every seat**, with **zero non-comparable cells anywhere** (NO_VERDICT - excluded 0, unmatched 0, all three arcs), so a difference between channels cannot be a difference in what was administered.
- Coverage: **16 of 16** sealed Core-16 seats, uncovered: **[]**, recomputed from the sealed constants, not asserted.

src: FP-M §3.1, §5.1 · FP-R §0.1 · CR §1

Kohlberg schema distributions — per seat, both coder legs, never averaged

Per-seat Kohlberg schema distribution — share of the seat's own judged kohberg-P cells assigned to each **PRIMARY** schema, by coder leg. The three schemas are descriptive structures and are never ordered; one panel per seat, grouped by substrate (row); both coder legs shown, never reduced to one code. Substrates never merged. Seats in declared roster order.



Per-seat Kohlberg primary-schema distributions, one panel per seat (18 panels), grouped into never-merged substrate rows. BOTH coder legs are drawn in every panel as paired bars per schema — leg A and leg B — never reduced to one code, never tie-broken, never averaged. Denominator: each seat's own 38 kohberg-P cells. The three schemas appear in sealed declared order and that order is NOT a scale — not stages, not levels, no developmental direction, register_without_structure is a count and is not drawn here. Both coders UNVALIDATED pre-Gate-M.

src: FP-R §2 Figure 3 · COMP §3.2 Figure 3b — *banked artifact, not re-rendered*

MFT — one profile shape, and a boundary that reverses

- On **both legs and all eight extension seats**, **purity** is the **smallest** of the six (leg A .125–.219; leg B .125–.156) and **care** is at or near the **largest** (leg A .875–1.000; leg B .844–.969).
- Every seat's profile there has the **same gross shape** — a statement about the instrument's **elicitation, not about seats relative to one another**.
- The **coder-boundary reversal**, over **108** seat × foundation comparisons (18 seat-blocks × 6 foundations): the **salient** rendering runs **66** leg-A-larger / 14 leg-B-larger / 28 tied; the **endorsed -only** rendering of the same cells runs **4 / 92 / 12**.
- The collapse **hides exactly the boundary the coders disagree on**. A **coder-pair property, not a seat property**; it does not tell us which coder is right.

src: FP-R §1.1, §1.4 · COMP §3.1

Kohlberg — three schemas, and they are NOT a scale

- `persona_interests` · `maintaining_norms` · `postconventional` · sealed declared order. **Not stages, not levels, no developmental direction**.
- register_without_structure** is a **COUNT, not a share**. A primary assignment is forced on every cell, so a flagged cell still carries one — the flag marks it **FORCED and low-confidence, not absent** — and flagged cells stay in the distributions by design.
- The **register-flag tally is scope-dependent and both forms are correct**: **40 (leg A) vs 7 (leg B) over 18 seat-blocks** (ext1 + confirmatory + the out-of-pool de-risk substrate), and **42 vs 7 over the three Core-16 arcs**. Never quoted without its scope.
- Leg A declares roughly **six times** as many of its own primary assignments to be forced: a difference in **annotation practice** that changes **no share** in any distribution.

src: FP-R §2, §2.3 · CR §3 · COMP §3.2

Schwartz-P — the absence IS the finding

No **per-seat 19-value priority profile exists anywhere in this program**. Not an omission and not a rendering choice — it is the state of the measurement.

- Two-stage and strictly ordered: a circumplex recovery gate on raw ratings **before** any priority vector. **No stage-2 output exists on any substrate**.
- 18 seat-gates** → **18 DECLINEs** on the results-comparison report's three $t = 0$ substrates, **plus 3 more on the coverage-completion arc** (every $t=0$ program total 21): **3 rating bundles per seat** (the geometric maximum under this design) against a sealed floor of **30 RESPONDENT** bundles. A design and budget fact, not a measurement outcome — and non-recovery is **NEVER** a model disposition.

QUARANTINED-DESCRIPTIVE until Gate M — not validity evidence

- Floor arc**, behind that verbatim banner: **PERMITTED at exactly 30 bundles** on all five seats (under stricter thresholds `harness_gate_passed = False` on all five — both readings reported, never merged, FP-R §3.4/§7.8), then `recovery_status = INDETERMINATE` on every seat — two of three sealed criteria have **no sealed implementation** and were not approximated. A statement about the **implementation**, not about any model. It ran at $t = 0.7$ (A-10) and is never numerically compared with any $t = 0$ surface.

src: FP-R §3.1, §3.4 · FP-M §5.5 · COMP §3.3.3

Channel reliability — what varies is the code, not the seat

Kohlberg channel agreement, **pooled over all 16 sealed Core-16 seats** and by arc; blocks share a column layout, **not comparability of rows**. Pooled denominators: **MFT $n = 512$** (16 × 32) · **kohlberg $n = 608$** (16 × 38). Agreement is a property of the **measurement instrument**, never of any model, and is **never per seat**.

kohlberg field	pooled, 16 seats (n = 608)	confirmatory (n = 190)	ext1 (n = 304)	ext2 (n = 114)
register_without_structure	0.933	0.884	0.977	0.895
primary	0.827	0.847	0.816	0.825
secondary	0.648	0.653	0.651	0.632
confidence	0.538	0.505	0.595	0.439

- The ordering **register > primary > secondary > confidence** holds on all three arcs **without exception** — three administrations, three disjoint rosters, three separate judge runs, one ordering. A single pooled column could be an artifact of one arc dominating the denominator; this is not. It also holds on a **fourth, out-of-pool substrate** (`dual-judge-p-e2e-v1`, $n = 190$) as a replication of the **ordering only**: substrate hygiene forbids pooling it, and none of its numbers is merged above.
- MFT, pooled (n = 512)**: **purity** 0.943 separates stably on every arc; the other five sit in a **0.660-0.785** band whose internal order is **not stable** — one separation, **not a ranking of six**, and never averaged into one MFT number.
- No subject seat is named against any figure on this panel**; seats are referred to only by arc, count, vendor family or min-max range.

src: CR §1, §2.1–§2.3 · COMP Tables 3.3.4a–b

Limitations an owner must carry

- Reliability, not validity** — the **load-bearing limitation**. Nothing here shows any code is correct. **Human expert gold (Gate M)** remains the **outstanding blocker for every validity claim in Family P**; it is user-owned and external, and is not advanced by one cell.
- Limitation 16, as annotated** — the **mft profile schema omits the confidence field** the **BINDING** specification requires. A missing field, not a wrong value: nothing computed is incorrect and no figure depends on mft confidence, because none was ever produced. The **evidence requirement** — **the part that makes a code auditable** — **IS met** on both P schemas. Ruled **DISCLOSE-not-repair**; no schema, prompt, rubric, pool or seal was touched. Read **"sealed"** as **BINDING** — the protocol section became binding at freeze but was never hash-sealed.
- The **judge-leg census predicate is completeness-blind**. `census_clean = identity_holds ^ unexpected == 0`; `terminal == scheduled` is not a term, so a **SHORT leg reads clean**. The compensating judge-leg guard has **ZERO** tests. It does not retroactively weaken the six legs — their final receipts show full and equal denominators — but the flag alone would not have caught a short leg.
- One coder pair**. Every pooled figure is a property of `claude-opus-5` × `gpt-5.6-sol` on this instrument; a different pair could produce a different channel profile. The A-4 four-coder arm cannot settle it: on the coverage-completion arc, five of six pairs carry `n_compared = 0`, and the one populated pair (7 kohberg / 8 MFT cells) includes neither production leg on the S7 confirmatory substrate all six pairs are populated (koh 7 / mft 8 cells each, production pair included — tiny n, but it exists; FP-R Table 4.1a). One of the six is a within-vendor pair, retained and always disclosed.
- Per-seat figures are not seat findings** — 32 and 38 cells per seat, so one cell moves a per-seat rate by 0.031 or 0.026, and no seat has a second administration reporting law permits comparing against. **No dispersion, no norming, no baselines**, and human-norm comparison is forbidden for P-KOH by construct design.
- Provenance and exposure**. These figures are read from a **banked, custody-certified** Stage-11 copy that cannot presently be re-rendered against the live study tree, and **exposure verification was WAIVED by the user** on the coverage-completion arc (CR §7.8 scope; FP-R §7.8 records the same-class waiver on the three *rendered* roster-ext seats).

src: FP-R §7 (16 as annotated) · CR §7 · COMP §3.2.3, §3.4

The dual-coder pair, and the disclosure it carries

- Two independent coders from different vendors, FIXED on the six-condition dossier** (§4.2 condition 7, MPD-178/D-163 — the pair and its condition-1-6 evidence were recorded in a decision entry before the first S7 cell; leg A `claude-opus-5` (Anthropic, hermetic Claude CLI OAuth, $c \leq 8$) and leg B `gpt-5.6-sol` (OpenAI, Codex CLI OAuth, $c = 1$ HARD). Both **\$0-cash OAuth**; realized **coder-leg cash \$0 on every substrate** (the ext2 arc's subject-side TRUE cash is \$0.86716843, correction disclosed in the assessment readout).
- The legs run **sequentially, never interleaved** (two guards race); caches are content-addressed and **coder-parameterised** on a fresh directory per run id.
- Both legs are always printed** — never reduced to one code, never tie-broken, never averaged.
- No leave-one-out in Family P**, deliberately. In-family exposure is 5 of 8 on the ext1 arc (3 Anthropic seats against leg A, 2 OpenAI against leg B) and 7 of 16 across the sealed **Core-16 roster** (3 Anthropic, 4 OpenAI) — never stated without its roster. Scope: **MFT-P/KOH-P cells only**.
- Disposition `retained_and_disclosed_never_excluded_never_ranked`. A **leniency or familiarity residual cannot be excluded and no arm of this design can detect it** — it is **CONTAINED by the never-rank law, not solved by it**.

src: COMP §3.2.1 · FP-M §4.2–§4.3 · CR §7.2

Census — with the resume record in front

- Two of the six executed coder legs reached final stats via ONE resume each**, both on the same exception class (a single unclassified `subprocess.TimeoutExpired`), both dispositioned **RESUME ONCE with sentinels archived rather than deleted**. Both retries succeeded — **transient, not systematic**. Neither abort is evidence against the pair.
- Six legs: **1,645 records per leg, 3,290 in total, every one census-clean**, with `identity_holds: true`, `rerun_targets: []` and `terminal_unmeasured_unexpected: 0` on all six; denominators equal leg-to-leg in every arc. The live proof travels with the predicate: the ext2 leg-A receipt reads `census_clean: true` at 206/210 with `aborted: true` — resumed once, finished 210/210, sentinel archived. **Seal vintage**: each S7 arc ran against its own pinned pool snapshot; the E-4 seal [2026-08-11, +16 pairs] postdates every record on this sheet.
- 16/16 coverage, uncovered: []** recomputed from the sealed constants; the three arcs are **disjoint by construction**, asserted at import.
- Resumed-leg process counters read **below** their census (`n_valid` counts this invocation's records; the census counts durable records). **The census is used everywhere**.
- Blinding across the three arcs: **1,604 prompts, 0 hits, 0 blocking**.

src: COMP §3.2.3 · FP-R §7.11

What none of this supports until Gate M

Does not. No ranking, no board, no composite, no tier, no ordering of seats, no claim about model quality, and **no confirmatory use** — reusing anything here that way **requires a new protocol**. It does not establish that any code is **CORRECT**. It does **not** cover schwartz-P, whose channels' reproducibility is simply **unknown**. It does **not** unblock Gate M, reduce what Gate M must do, or **substitute for human expert gold** in any degree.

Reliability is not validity, and only reliability was measured. Agreement rates are a property of the measurement instrument, never of any model. **NEVER RANKED stays absolute and by construction** — even a passed Gate M would not create a direction on these constructs. **\$0 session**: every figure is read from a banked artifact, no model or API call was made. **No benchmark item text appears anywhere on this sheet** — **ids, counts and aggregates only**; seat names appear only as figure panel labels in declared roster order, and no per-seat agreement or kappa quantity is attached to any seat anywhere in this program.

src: CR §8 · COMP §3.4, §4b.7 · FP-R A-1 render law