

Publish task results. Do not rank the models.

CEI moral psychology benchmark

Research lead readout · 4 September 2026

Why: the highest saved score changes by task, and saved ranges overlap on both comparison tests.

Share the task results with their limits. Size and release slides are early clues, not final answers.

No model has the highest saved score on all eight tasks

MAIN RESULTS

Highest saved score for each task, using the same five models

Task	Highest saved score	Score type	Model
MFQ agreement	.884	Agreement score	GPT-5.4 mini
Vignette agreement	.920	Agreement score	GPT-5.4
MFQ compare	.550	Accuracy	Haiku + Qwen tie†
Vignette compare	.625	Accuracy	Haiku + GPT-5.4 tie†
UniMoral action	.668	Accuracy	Claude Haiku 4.5
UniMoral typology	.651	Accuracy	Claude Opus 4.8
UniMoral factor	.609	Accuracy	Claude Haiku 4.5
UniMoral consequence	.152	Text-match score	Claude Opus 4.8

† The uncertainty ranges overlap on the two comparison tasks. We cannot tell who leads.


Each task uses its own score. Do not compare numbers from different rows.

Saved ranges overlap for every model pair in both tests

MAIN RESULTS

Full primary panels: MFQ = 8 models × 20 questions; Vignette = 10 × 24.

Each bar = the median full width of saved 95% intervals across available models.



Direct evidence: saved ranges overlap for all 28 MFQ and all 45 Vignette model pairs.


Action: recover per-question outcomes and check scoring.

Then compare models, run human review, and add questions only if still unclear.

Only 4 of 12 model-and-task cases rise twice

EXPLORATORY

Qwen, Gemma, and Llama × four tasks; model versions differ in more than size




Each bar counts model-family × task cases (12 total).

Size is one clue, not a general rule.

Each case follows one model family on one task. Different score types stay separate.

Gemma scores rise on one task and fall on another

EXPLORATORY



The larger named variants do not score higher on both tasks.

Each point is one stored task score. Its uncertainty range and raw run archive are unavailable.

Later model versions score higher on some tasks and lower on others


EXPLORATORY

Score change on three UniMoral tasks

Qwen: 2024-Q3 · Qwen2.5 7B Instruct (7.61B) → 2026-Q1 · Qwen3.5 9B (9B)

DeepSeek: 2025-Q1 · V3-0324 (671B main, 37B active) → 2026-Q2 · V4 Flash (284B main, 13B active)

main = published main-model parameters (auxiliary/MTP excluded); active = parameters used per token



All endpoint runs were evaluated on 28–29 May 2026.

Release period is model metadata, not a progress timeline; 28–29 May is the evaluation time.

Exploratory: saved uncertainty is unavailable. Consequence uses a different score and is not shown.

We did not exactly repeat any paper's experiment

The papers help explain our questions. Their scores cannot be compared with ours.

0 of 4
papers repeated
exactly

Paper	Plain-language question	How close is our test?
MoralBench	Do model choices match human ratings?	Similar question
UniMoral	Can models predict choices, moral categories, influences, and what happens next?	Some similar tasks
MoReBench	Does reasoning cover expert criteria?	Different scoring
MoralLens	Do reasons change when a model explains before or after choosing?	Different scoring

Use the papers to explain our questions—not to compare scores.

Our UniMoral run does not test prompt hints, language differences, or where the stories came from.

Share task results now. Strengthen the test next.

DECISION

Better evidence makes the current results more useful

When	Action	Reason
Now	Share each task result with its limits	The saved results answer one task at a time
Next	Recover per-question outcomes; check scoring and labels	This lets us compare models directly
Then	Have people review the test	A benchmark score does not prove the test matches human judgment
Only if still unclear	Add more comparison questions	New questions help after the scoring works correctly

Best research next move

Recover per-question outcomes. Compare models. Then have people review the test.