

MoralBench + UniMoral

No model has the highest saved score on all 8 tasks.

MoralBench comparison

For every model pair, the two models' 95% score ranges overlap: 28 of 28 MFQ pairs and 45 of 45 vignette pairs.

UniMoral model size and release date

Scores move up and down across tasks. Bigger or newer is not consistently higher here.

MoralBench + UniMoral: no model has the top saved score on all 8 tasks

CEI local result. Same five models; each row keeps its own score type.

Task	Highest saved score	Score type	Model
MFQ agreement Track benchmark human ratings of value statements	.884	Agreement score	GPT-5.4 mini
Vignette agreement Track benchmark human ratings of moral stories	.920	Agreement score	GPT-5.4
MFQ compare Pick the higher-rated value statement	.550	Accuracy	Haiku + Qwen tie†
Vignette compare Pick the higher-rated moral story	.625	Accuracy	Haiku + GPT-5.4 tie†
UniMoral action Match one recorded human choice	.668	Accuracy	Claude Haiku 4.5
UniMoral typology Classify the recorded choice into one of four types	.651	Accuracy	Claude Opus 4.8
UniMoral factor Pick a top-rated factor for the recorded choice	.609	Accuracy	Claude Haiku 4.5
UniMoral consequence Write what could happen next	.152	Text-match score	Claude Opus 4.8

† Every 95% range overlaps on both MoralBench comparison tasks. The apparent leaders are unresolved.

Each task uses its own score. Do not compare numbers from different rows.

MoralBench: each pair of models has overlapping 95% score ranges

CEI local result. MFQ uses 8 models and 20 questions. Vignettes use 10 models and 24 questions.

MFQ value statements

28 / 28

model pairs have overlapping 95% score ranges

Moral vignettes

45 / 45


model pairs have overlapping 95% score ranges

These saved results do not separate a leader.

The ranges are not a direct model-pair test. Question-level scores and stronger uncertainty estimates are unavailable.

UniMoral: only 4 of 12 family and task comparisons rise twice

Exploratory CEI result. Each comparison follows one model family on one task across three sizes.




Each bar counts one model family on one task across three saved sizes.

4 rise twice · 7 change direction · 1 falls twice

This is exploratory. Model versions also differ in training and release, and no confidence ranges were saved.

UniMoral Gemma: factor rises while typology falls

Exploratory CEI result. Gemma 3 at 4B, 12B, and 27B; the vertical axis is zoomed.



4B to 27B: factor accuracy $+0.034$ · moral typology accuracy -0.027

Each point is one saved task score. No uncertainty range or raw run archive is available.


UniMoral releases: Qwen rises on 3 tasks, DeepSeek on 2

Exploratory CEI result. Later score minus earlier score; consequence uses text match.

Qwen: Qwen2.5 7B Instruct · 2024-09-19 · 7.61B total to Qwen3.5 9B · 2026-02-27 · 9B total

DeepSeek: V3-0324 · 2025-03-24 · 671B main / 37B active to V4 Flash · 2026-04-22 · 284B main / 13B active

DeepSeek labels show published main-model B and active B used per token.



Across all 4 tasks: Qwen is higher on 3 and lower on 1. DeepSeek is higher on 2 and lower on 2.

Separate metric · Consequence text match (METEOR): Qwen -0.027 · DeepSeek $+0.019$

All scores were tested 28–29 May 2026. Release date describes the model, not the test date. No uncertainty ranges were saved.

How the local tests relate to five papers

MoralBench, UniMoral, MoReBench, MoralLens, and Value Kaleidoscope

Paper	What it asks	Closest local evidence
MoralBench	Do model choices match human ratings?	Related tasks, different setup
UniMoral	Can models predict actions, types, factors, and consequences?	Narrower versions of four tasks
MoReBench	Does reasoning meet expert criteria?	Keyword check only
MoralLens	Does reasoning order change reasons and choices?	Keyword check only
Value Kaleidoscope	How do value tasks change with model size?	Different models and scoring

None of the five is reproduced exactly. Paper scores and local scores stay separate.

MoralBench and UniMoral: what we can report now

Report the 8 task results now. Do not publish one overall model rank.

When	Action	Reason
Now	Share each task result with its limits	The saved results answer one task at a time
Next	Restore each model's answer and score for every question	This lets us compare models on the same questions
Then	Check scoring and labels; have people review the test	A benchmark score alone does not prove the test matches human judgment
Only if still unclear	Add more comparison questions	New questions help after the scoring works correctly

Next: restore answers and scores, check scoring, compare models, and get human review.