

Publish task results. Do not rank the models.

CEI moral psychology benchmark

Research lead readout · 4 September 2026

Why: the top model changes by task, and two comparison tests still cannot give a reliable order.

Share the task results with their limits. Size and release slides are early clues, not final answers.

No model leads all eight tasks

MAIN RESULTS

Highest stored result for each task, using the same five models


Task	Highest stored result	Score type	Model
MFQ agreement	.884	Agreement score	GPT-5.4 mini
Vignette agreement	.920	Agreement score	GPT-5.4
MFQ compare	.550	Accuracy	Haiku + Qwen tie†
Vignette compare	.625	Accuracy	Haiku + GPT-5.4 tie†
UniMoral action	.668	Accuracy	Claude Haiku 4.5
UniMoral typology	.651	Accuracy	Claude Opus 4.8
UniMoral factor	.609	Accuracy	Claude Haiku 4.5
UniMoral consequence	.152	Text-match score	Claude Opus 4.8

† The uncertainty ranges overlap on the two comparison tasks. We cannot tell who leads.

Each task uses its own score. Do not compare numbers from different rows.

Two comparison tests cannot rank models yet

Wider bar = less sure. MFQ has 20 questions per model; Vignette has 24.




Action: compare every model on the same questions.

If the order is still unclear, add more comparison questions.

Only 4 of 12 model-and-task cases rise twice

EXPLORATORY

Qwen, Gemma, and Llama × four tasks; model versions differ in more than size




Size is one clue, not a general rule.

Each case follows one model family on one task. Different score types stay separate.

The same Gemma models move in opposite directions

EXPLORATORY



The larger named variants do not score higher on both tasks.

Each point is one stored task score. Its uncertainty range and original run are unavailable.

Later model versions score higher on some tasks and lower on others


EXPLORATORY

Score change on three UniMoral tasks

Qwen: Qwen2.5 7B Instruct (7.61B) to Qwen3.5 9B (9B)

DeepSeek: V3-0324 (671B main, 37B active) to V4 Flash (284B main, 13B active)

main = all parameters; active = parameters used for each token



All endpoint runs were evaluated on 28–29 May 2026.

Release period describes the model. It is not a progress timeline.

The consequence task uses a different score, so it is not shown here.

We did not exactly repeat any paper's experiment

The papers help explain our questions. Their scores cannot be compared with ours.

0

**papers repeated
exactly**

Paper	Plain-language question	How close is our test?
MoralBench	Do model choices match human ratings?	Similar question
UniMoral	Can models predict choices, moral categories, influences, and what happens next?	Some similar tasks
MoReBench	Does reasoning cover expert criteria?	Different scoring
MoralLens	Do reasons change when a model explains before or after choosing?	Different scoring

Use the papers to explain our questions—not to compare scores.

Our UniMoral run does not test prompt hints, language differences, or where the stories came from.

Make the test stronger before adding more models

DECISION

Better evidence makes the current results more useful

When	Action	Reason
Now	Share each task result with its limits	The saved results answer one task at a time
Next	Compare the same questions; check that answers were read and labeled correctly	This tackles the biggest gaps
Then	Have people review the test	A benchmark score does not prove the test matches human judgment
Only if still unclear	Add more comparison questions	New questions help after the scoring works correctly

Best next move

Compare models on the same questions. Then have people review the test.