Publish task results.
Do not rank the models.

CEI moral psychology benchmark

Research lead readout · 4 September 2026

Why: the highest saved score changes by task, and saved ranges overlap on both comparison tests.

Share the task results with their limits. Size and release slides are early clues, not final answers.

No model has the highest saved score on all eight tasks

MAIN RESULTS

Highest saved score for each task, using the same five models

Task	Highest saved score	Score type	Model
MFQ agreement Track benchmark human ratings of value statements	.884	Agreement score	**GPT-5.4 mini**
Vignette agreement Track benchmark human ratings of moral stories	.920	Agreement score	**GPT-5.4**
MFQ compare Pick the higher-rated value statement	.550	Accuracy	**Haiku + Qwen tie†**
Vignette compare Pick the higher-rated moral story	.625	Accuracy	**Haiku + GPT-5.4 tie†**
UniMoral action Match one recorded human choice	.668	Accuracy	**Claude Haiku 4.5**
UniMoral typology Classify the recorded choice into one of four types	.651	Accuracy	**Claude Opus 4.8**
UniMoral factor Pick a top-rated factor for the recorded choice	.609	Accuracy	**Claude Haiku 4.5**
UniMoral consequence Write what could happen next	.152	Text-match score	**Claude Opus 4.8**

† Marginal ranges overlap on both comparison tasks; paired question results are unavailable, so no leader is resolved.

Each task uses its own score. Do not compare numbers from different rows.

Every model pair has overlapping score ranges on both tests

MAIN RESULTS

Full primary panels: MFQ = 8 models × 20 questions; Vignette = 10 × 24.

Each bar = the median full width of saved 95% intervals across available models.

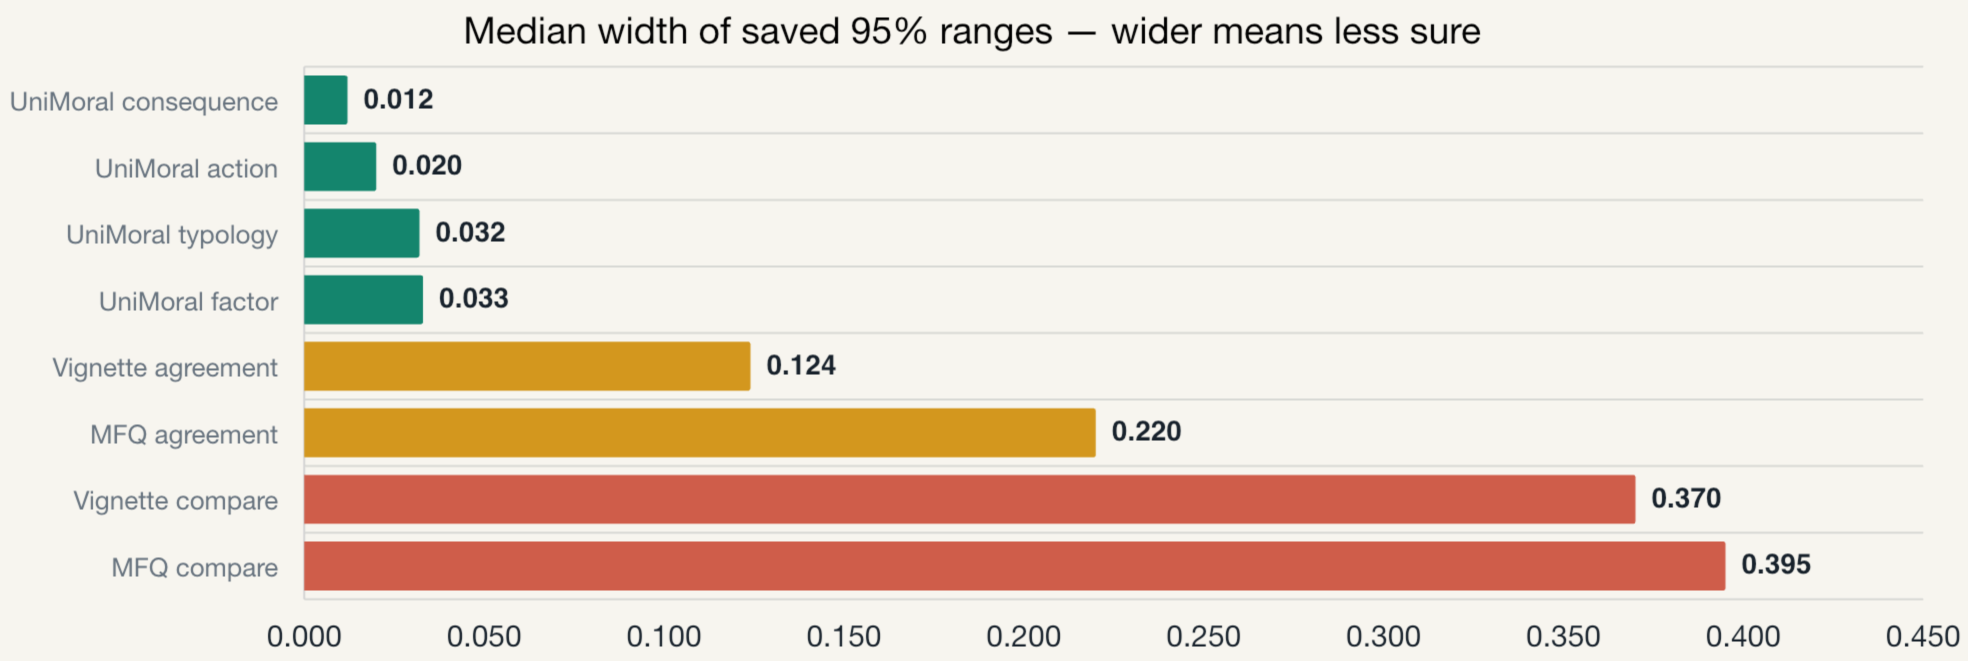

Marginal ranges overlap for all 28 MFQ and all 45 Vignette model pairs.

These ranges look at one model at a time; paired question results are unavailable.

Next: restore every model's answer and score for each question.

Check scoring and labels. Then compare models and have people review the test.

Only 4 of 12 model-and-task cases rise twice

EXPLORATORY

Qwen, Gemma, and Llama × four tasks; model versions differ in more than size

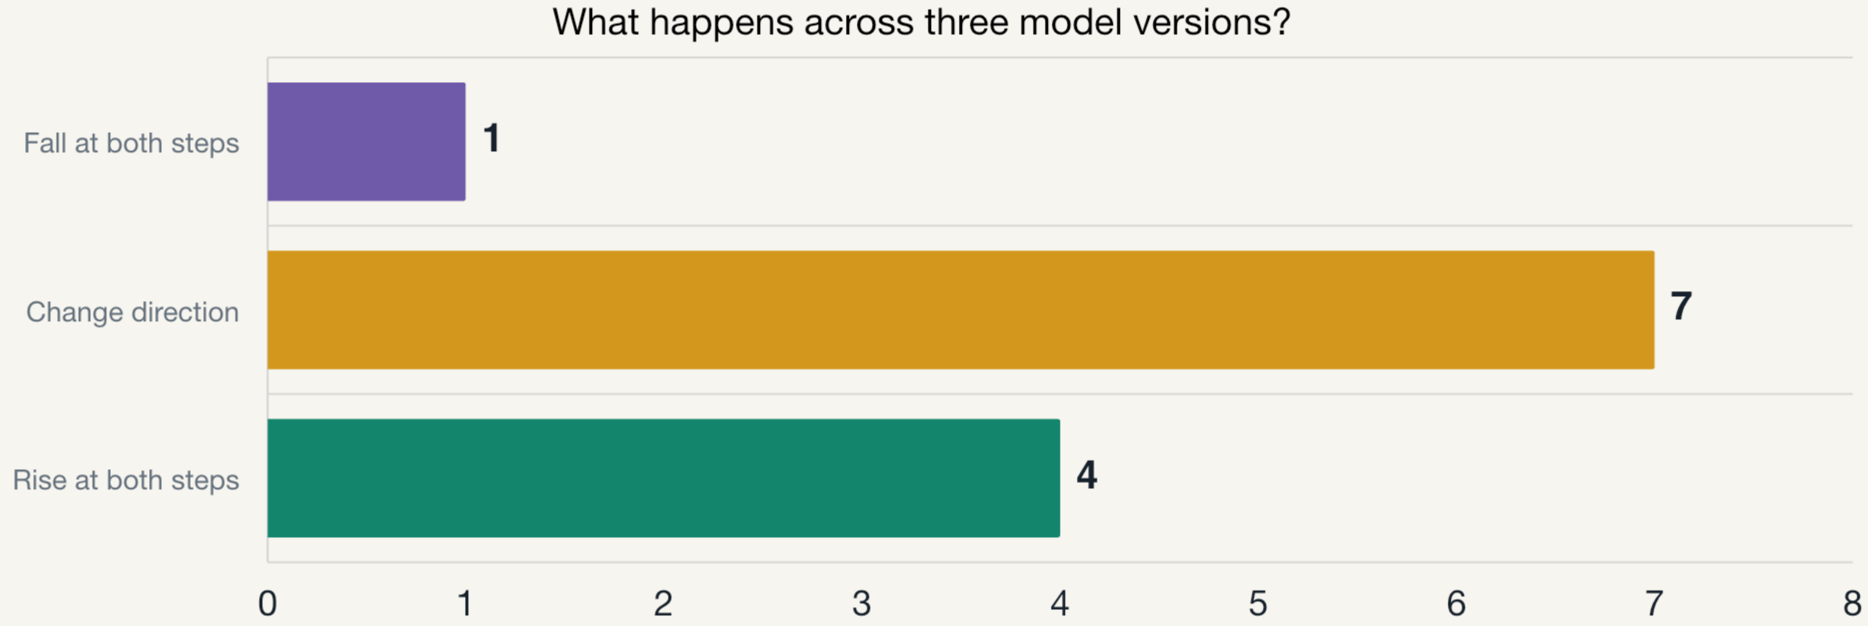

Each bar counts model-family × task cases (12 total).

Size is one clue, not a general rule.

Each case follows one model family on one task. Different score types stay separate.

Gemma scores rise on one task and fall on another

EXPLORATORY

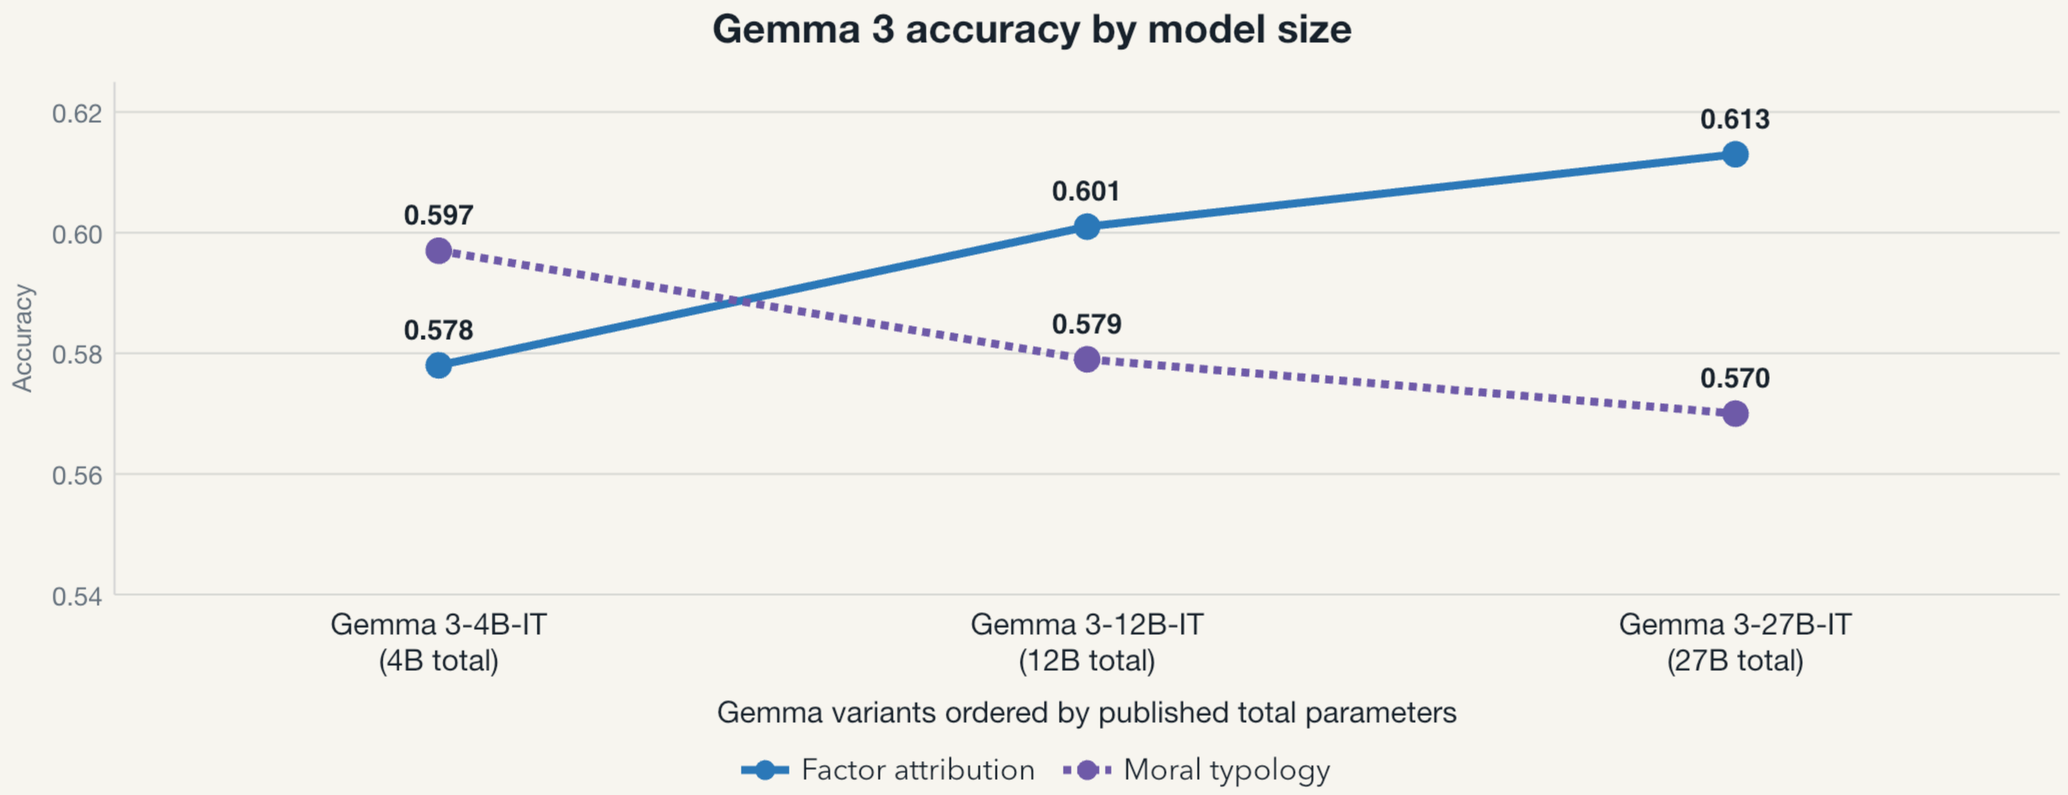

The larger named variants do not score higher on both tasks.

Each point is one stored task score. Its uncertainty range and raw run archive are unavailable.

Later model versions score higher on some tasks and lower on others

EXPLORATORY

Score change on three UniMoral tasks

Qwen: 2024-Q3 · Qwen2.5 7B Instruct (7.61B) → 2026-Q1 · Qwen3.5 9B (9B)

DeepSeek: 2025-Q1 · V3-0324 (671B main, 37B active) → 2026-Q2 · V4 Flash (284B main, 13B active)

main = published main-model parameters (auxiliary/MTP excluded); active = parameters used per token

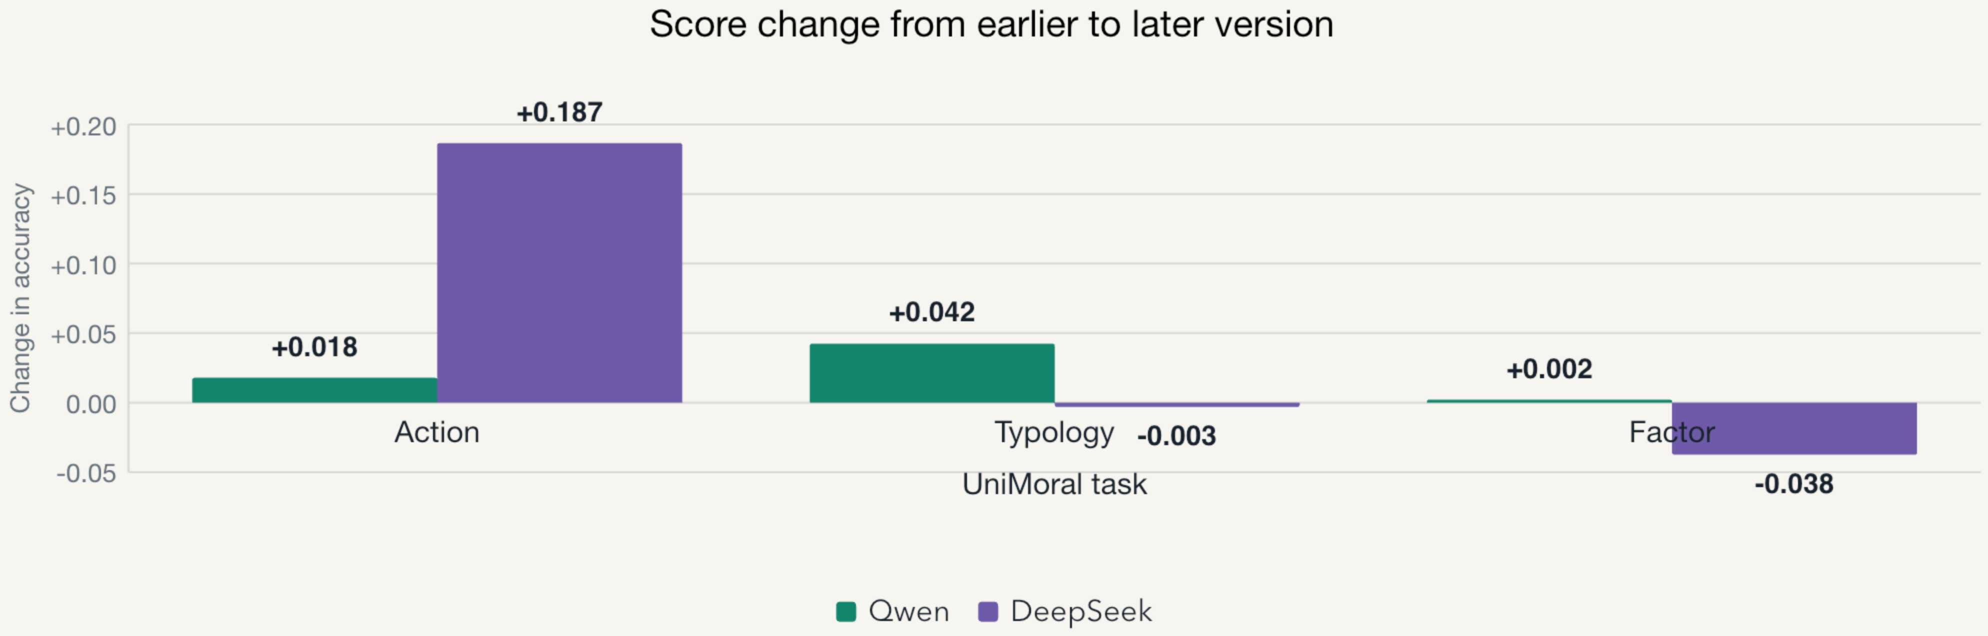

All endpoint runs were evaluated on 28–29 May 2026.

Release period is model metadata, not a progress timeline; 28–29 May is the evaluation time.

Exploratory: saved uncertainty is unavailable. Consequence uses a different score and is not shown.

We did not exactly repeat any paper's experiment

PAPER REVIEW

The papers help explain our questions. Their scores cannot be compared with ours.

0 of 4

papers repeated exactly

Paper	Plain-language question	How close is our test?
MoralBench	Do model choices match human ratings?	**Similar question**
UniMoral	Can models predict choices, moral categories, influences, and what happens next?	**Some similar tasks**
MoReBench	Does reasoning cover expert criteria?	**Related question, different test**
MoralLens	Do reasons change when a model explains before or after choosing?	**Related question, different test**

Use the papers to explain our questions—not to compare scores.

Our UniMoral run does not test prompt hints, language differences, or where the stories came from.

Share task results now. Strengthen the test next.

DECISION

Better evidence makes the current results more useful

When	Action	Reason
Now	Share each task result with its limits	The saved results answer one task at a time
Next	Restore each model's answer and score for every question	This lets us compare models on the same questions
Then	Check scoring and labels; have people review the test	A benchmark score alone does not prove the test matches human judgment
Only if still unclear	Add more comparison questions	New questions help after the scoring works correctly

Best research next move

Restore answers and scores. Check scoring and labels. Compare models. Then have people review the test.